

Stochastic Dual Coordinate Ascent Methods for Regularized Loss Minimization

Shai Shalev-Shwartz

*Benin school of Computer Science and Engineering
The Hebrew University
Jerusalem, 91904, Israel*

SHAIS@CS.HUJI.AC.IL

Tong Zhang

*Department of Statistics
Rutgers University
Piscataway, NJ, 08854, USA*

TZHANG@STAT.RUTGERS.EDU

Editor: Leon Bottou

Abstract

Stochastic Gradient Descent (SGD) has become popular for solving large scale supervised machine learning optimization problems such as SVM, due to their strong theoretical guarantees. While the closely related Dual Coordinate Ascent (DCA) method has been implemented in various software packages, it has so far lacked good convergence analysis. This paper presents a new analysis of Stochastic Dual Coordinate Ascent (SDCA) showing that this class of methods enjoy strong theoretical guarantees that are comparable or better than SGD. This analysis justifies the effectiveness of SDCA for practical applications.

Keywords: stochastic dual coordinate ascent, optimization, computational complexity, regularized loss minimization, support vector machines, ridge regression, logistic regression

1. Introduction

We consider the following generic optimization problem associated with regularized loss minimization of linear predictors: Let x_1, \dots, x_n be vectors in \mathbb{R}^d , let ϕ_1, \dots, ϕ_n be a sequence of scalar convex functions, and let $\lambda > 0$ be a regularization parameter. Our goal is to solve $\min_{w \in \mathbb{R}^d} P(w)$ where¹

$$P(w) = \left[\frac{1}{n} \sum_{i=1}^n \phi_i(w^\top x_i) + \frac{\lambda}{2} \|w\|^2 \right]. \quad (1)$$

For example, given labels y_1, \dots, y_n in $\{\pm 1\}$, the SVM problem (with linear kernels and no bias term) is obtained by setting $\phi_i(a) = \max\{0, 1 - y_i a\}$. Regularized logistic regression is obtained by setting $\phi_i(a) = \log(1 + \exp(-y_i a))$. Regression problems also fall into the above. For example, ridge regression is obtained by setting $\phi_i(a) = (a - y_i)^2$, regression with the absolute-value is obtained by setting $\phi_i(a) = |a - y_i|$, and support vector regression is obtained by setting $\phi_i(a) = \max\{0, |a - y_i| - \nu\}$, for some predefined insensitivity parameter $\nu > 0$.

1. Throughout this paper, we only consider the ℓ_2 -norm.

Let w^* be the optimum of (1). We say that a solution w is ϵ_P -sub-optimal if $P(w) - P(w^*) \leq \epsilon_P$. We analyze the runtime of optimization procedures as a function of the time required to find an ϵ_P -sub-optimal solution.

A simple approach for solving SVM is stochastic gradient descent (SGD) (Robbins and Monro, 1951; Murata, 1998; Cun and Bottou, 2004; Zhang, 2004; Bottou and Bousquet, 2008; Shalev-Shwartz et al., 2007). SGD finds an ϵ_P -sub-optimal solution in time $\tilde{O}(1/(\lambda\epsilon_P))$. This runtime does not depend on n and therefore is favorable when n is very large. However, the SGD approach has several disadvantages. It does not have a clear stopping criterion; it tends to be too aggressive at the beginning of the optimization process, especially when λ is very small; while SGD reaches a moderate accuracy quite fast, its convergence becomes rather slow when we are interested in more accurate solutions.

An alternative approach is dual coordinate ascent (DCA), which solves a *dual* problem of (1). Specifically, for each i let $\phi_i^* : \mathbb{R} \rightarrow \mathbb{R}$ be the convex conjugate of ϕ_i , namely, $\phi_i^*(u) = \max_z (zu - \phi_i(z))$. The dual problem is

$$\max_{\alpha \in \mathbb{R}^n} D(\alpha) \quad \text{where} \quad D(\alpha) = \left[\frac{1}{n} \sum_{i=1}^n -\phi_i^*(-\alpha_i) - \frac{\lambda}{2} \left\| \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i x_i \right\|^2 \right]. \quad (2)$$

The dual objective in (2) has a different dual variable associated with each example in the training set. At each iteration of DCA, the dual objective is optimized with respect to a single dual variable, while the rest of the dual variables are kept in tact.

If we define

$$w(\alpha) = \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i x_i, \quad (3)$$

then it is known that $w(\alpha^*) = w^*$, where α^* is an optimal solution of (2). It is also known that $P(w^*) = D(\alpha^*)$ which immediately implies that for all w and α , we have $P(w) \geq D(\alpha)$, and hence the duality gap defined as

$$P(w(\alpha)) - D(\alpha)$$

can be regarded as an upper bound of the primal sub-optimality $P(w(\alpha)) - P(w^*)$.

We focus on a *stochastic* version of DCA, abbreviated by SDCA, in which at each round we choose which dual coordinate to optimize uniformly at random. The purpose of this paper is to develop theoretical understanding of the convergence of the duality gap for SDCA.

We analyze SDCA either for L -Lipschitz loss functions or for $(1/\gamma)$ -smooth loss functions, which are defined as follows. Throughout the paper, we will use $\phi'(a)$ to denote a sub-gradient of a convex function $\phi(\cdot)$, and use $\partial\phi(a)$ to denote its sub-differential.

Definition 1 A function $\phi_i : \mathbb{R} \rightarrow \mathbb{R}$ is L -Lipschitz if for all $a, b \in \mathbb{R}$, we have

$$|\phi_i(a) - \phi_i(b)| \leq L|a - b|.$$

A function $\phi_i : \mathbb{R} \rightarrow \mathbb{R}$ is $(1/\gamma)$ -smooth if it is differentiable and its derivative is $(1/\gamma)$ -Lipschitz. An equivalent condition is that for all $a, b \in \mathbb{R}$, we have

$$\phi_i(a) \leq \phi_i(b) + \phi_i'(b)(a - b) + \frac{1}{2\gamma}(a - b)^2,$$

where ϕ_i' is the derivative of ϕ_i .

It is well-known that if $\phi_i(a)$ is $(1/\gamma)$ -smooth, then $\phi_i^*(u)$ is γ strongly convex: for all $u, v \in \mathbb{R}$ and $s \in [0, 1]$:

$$-\phi_i^*(su + (1-s)v) \geq -s\phi_i^*(u) - (1-s)\phi_i^*(v) + \frac{\gamma s(1-s)}{2}(u-v)^2.$$

Our main findings are: in order to achieve a duality gap of ϵ ,

- For L -Lipschitz loss functions, we obtain the rate of $\tilde{O}(n + L^2/(\lambda\epsilon))$.
- For $(1/\gamma)$ -smooth loss functions, we obtain the rate of $\tilde{O}((n + 1/(\lambda\gamma)) \log(1/\epsilon))$.
- For loss functions which are almost everywhere smooth (such as the hinge-loss), we can obtain rate better than the above rate for Lipschitz loss. See Section 5 for a precise statement.

2. Related Work

DCA methods are related to decomposition methods (Platt, 1998; Joachims, 1998). While several experiments have shown that decomposition methods are inferior to SGD for large scale SVM (Shalev-Shwartz et al., 2007; Bottou and Bousquet, 2008), Hsieh et al. (2008) recently argued that SDCA outperform the SGD approach in some regimes. For example, this occurs when we need relatively high solution accuracy so that either SGD or SDCA has to be run for more than a few passes over the data.

However, our theoretical understanding of SDCA is not satisfying. Several authors (e.g., Mangasarian and Musicant, 1999; Hsieh et al., 2008) proved a linear convergence rate for solving SVM with DCA (not necessarily stochastic). The basic technique is to adapt the linear convergence of coordinate ascent that was established by Luo and Tseng (1992). The linear convergence means that it achieves a rate of $(1 - \nu)^k$ after k passes over the data, where $\nu > 0$. This convergence result tells us that after an unspecified number of iterations, the algorithm converges faster to the optimal solution than SGD.

However, there are two problems with this analysis. First, the linear convergence parameter, ν , may be very close to zero and the initial unspecified number of iterations might be very large. In fact, while the result of Luo and Tseng (1992) does not explicitly specify ν , an examine of their proof shows that ν is proportional to the smallest nonzero eigenvalue of $X^\top X$, where X is the $n \times d$ data matrix with its i -th row be the i -th data point x_i . For example if two data points $x_i \neq x_j$ becomes closer and closer, then $\nu \rightarrow 0$. This dependency is problematic in the data laden domain, and we note that such a dependency does not occur in the analysis of SGD.

Second, the analysis only deals with the sub-optimality of the *dual* objective, while our real goal is to bound the sub-optimality of the *primal* objective. Given a dual solution $\alpha \in \mathbb{R}^n$ its corresponding primal solution is $w(\alpha)$ (see (3)). The problem is that even if α is ϵ_D -sub-optimal in the dual, for some small ϵ_D , the primal solution $w(\alpha)$ might be far from being optimal. For SVM, (Hush et al., 2006, Theorem 2) showed that in order to obtain a primal ϵ_P -sub-optimal solution, we need a dual ϵ_D -sub-optimal solution with $\epsilon_D = O(\lambda\epsilon_P^2)$; therefore a convergence result for dual solution can only translate into a primal convergence result with worse convergence rate. Such a treatment is unsatisfactory, and this is what we will avoid in the current paper.

Some analyses of stochastic coordinate ascent provide solutions to the first problem mentioned above. For example, Collins et al. (2008) analyzed an exponentiated gradient dual coordinate ascent algorithm. The algorithm analyzed there (exponentiated gradient) is different from the standard DCA algorithm which we consider here, and the proof techniques are quite different. Consequently their results are not directly comparable to results we obtain in this paper. Nevertheless we note that for SVM, their analysis shows a convergence rate of $O(n/\epsilon_D)$ in order to achieve ϵ_D -sub-optimality (on the dual) while our analysis shows a convergence of $O(n \log \log n + 1/\lambda\epsilon)$ to achieve ϵ duality gap; for logistic regression, their analysis shows a convergence rate of $O((n + 1/\lambda) \log(1/\epsilon_D))$ in order to achieve ϵ_D -sub-optimality on the dual while our analysis shows a convergence of $O((n + 1/\lambda) \log(1/\epsilon))$ to achieve ϵ duality gap.

In addition, Shalev-Shwartz and Tewari (2009), and later Nesterov (2012) have analyzed randomized versions of coordinate descent for unconstrained and constrained minimization of smooth convex functions. Hsieh et al. (2008, Theorem 4) applied these results to the dual SVM formulation. However, the resulting convergence rate is $O(n/\epsilon_D)$ which is, as mentioned before, inferior to the results we obtain here. Furthermore, neither of these analyses can be applied to logistic regression due to their reliance on the smoothness of the dual objective function which is not satisfied for the dual formulation of logistic regression. We shall also point out again that all of these bounds are for the dual sub-optimality, while as mentioned before, we are interested in the primal sub-optimality.

In this paper we derive new bounds on the duality gap (hence, they also imply bounds on the primal sub-optimality) of SDCA. These bounds are superior to earlier results, and our analysis only holds for randomized (stochastic) dual coordinate ascent. As we will see from our experiments, randomization is important in practice. In fact, the practical convergence behavior of (non-stochastic) cyclic dual coordinate ascent (even with a random ordering of the data) can be slower than our theoretical bounds for SDCA, and thus cyclic DCA is inferior to SDCA. In this regard, we note that some of the earlier analysis such as Luo and Tseng (1992) can be applied both to stochastic and to cyclic dual coordinate ascent methods with similar results. This means that their analysis, which can be no better than the behavior of cyclic dual coordinate ascent, is inferior to our analysis.

Recently, Lacoste-Julien et al. (2012) derived a stochastic coordinate ascent for structural SVM based on the Frank-Wolfe algorithm. Specifying one variant of their algorithm to binary classification with the hinge loss, yields the SDCA algorithm for the hinge-loss. The rate of convergence Lacoste-Julien et al. (2012) derived for their algorithm is the same as the rate we derive for SDCA with a Lipschitz loss function.

Another relevant approach is the Stochastic Average Gradient (SAG), that has recently been analyzed in Le Roux et al. (2012). There, a convergence rate of $\tilde{O}(n \log(1/\epsilon))$ rate is shown, for the case of smooth losses, assuming that $n \geq \frac{8}{\lambda\gamma}$. This matches our guarantee in the regime $n \geq \frac{8}{\lambda\gamma}$.

The following table summarizes our results in comparison to previous analyses. Note that for SDCA with Lipschitz loss, we observe a faster practical convergence rate, which is explained with our refined analysis in Section 5.

Lipschitz loss		
Algorithm	type of convergence	rate
SGD	primal	$\tilde{O}(\frac{1}{\lambda\epsilon})$
online EG (Collins et al., 2008) (for SVM)	dual	$\tilde{O}(\frac{n}{\epsilon})$
Stochastic Frank-Wolfe (Lacoste-Julien et al., 2012)	primal-dual	$\tilde{O}(n + \frac{1}{\lambda\epsilon})$
SDCA	primal-dual	$\tilde{O}(n + \frac{1}{\lambda\epsilon})$ or faster

Smooth loss		
Algorithm	type of convergence	rate
SGD	primal	$\tilde{O}(\frac{1}{\lambda\epsilon})$
online EG (Collins et al., 2008) (for logistic regression)	dual	$\tilde{O}((n + \frac{1}{\lambda}) \log \frac{1}{\epsilon})$
SAG (Le Roux et al., 2012) (assuming $n \geq \frac{8}{\lambda\gamma}$)	primal	$\tilde{O}((n + \frac{1}{\lambda}) \log \frac{1}{\epsilon})$
SDCA	primal-dual	$\tilde{O}((n + \frac{1}{\lambda}) \log \frac{1}{\epsilon})$

3. Basic Results

The generic algorithm we analyze is described below. In the pseudo-code, the parameter T indicates the number of iterations while the parameter T_0 can be chosen to be a number between 1 to T . Based on our analysis, a good choice of T_0 is to be $T/2$. In practice, however, the parameters T and T_0 are not required as one can evaluate the duality gap and terminate when it is sufficiently small.

Procedure SDCA($\alpha^{(0)}$)
Let $w^{(0)} = w(\alpha^{(0)})$ Iterate: for $t = 1, 2, \dots, T$: Randomly pick i Find $\Delta\alpha_i$ to maximize $-\phi_i^*(-(\alpha_i^{(t-1)} + \Delta\alpha_i)) - \frac{\lambda n}{2} \ w^{(t-1)} + (\lambda n)^{-1} \Delta\alpha_i x_i\ ^2$ $\alpha^{(t)} \leftarrow \alpha^{(t-1)} + \Delta\alpha_i e_i$ $w^{(t)} \leftarrow w^{(t-1)} + (\lambda n)^{-1} \Delta\alpha_i x_i$ Output (Averaging option): Let $\bar{\alpha} = \frac{1}{T-T_0} \sum_{i=T_0+1}^T \alpha^{(i-1)}$ Let $\bar{w} = w(\bar{\alpha}) = \frac{1}{T-T_0} \sum_{i=T_0+1}^T w^{(i-1)}$ return \bar{w} Output (Random option): Let $\bar{\alpha} = \alpha^{(t)}$ and $\bar{w} = w^{(t)}$ for some random $t \in T_0 + 1, \dots, T$ return \bar{w}

We analyze the algorithm based on different assumptions on the loss functions. To simplify the statements of our theorems, we always assume the following:

1. For all i , $\|x_i\| \leq 1$
2. For all i and a , $\phi_i(a) \geq 0$

3. For all i , $\phi_i(0) \leq 1$

Theorem 2 Consider Procedure SDCA with $\alpha^{(0)} = 0$. Assume that ϕ_i is L -Lipschitz for all i . To obtain a duality gap of $\mathbb{E}[P(\bar{w}) - D(\bar{\alpha})] \leq \epsilon_P$, it suffices to have a total number of iterations of

$$T \geq T_0 + n + \frac{4L^2}{\lambda\epsilon_P} \geq \max(0, \lceil n \log(0.5\lambda n L^{-2}) \rceil) + n + \frac{20L^2}{\lambda\epsilon_P}.$$

Moreover, when $t \geq T_0$, we have dual sub-optimality bound of $\mathbb{E}[D(\alpha^*) - D(\alpha^{(t)})] \leq \epsilon_P/2$.

Remark 3 If we choose the average version, we may simply take $T = 2T_0$. Moreover, we note that Theorem 2 holds for both averaging or for choosing w at random from $\{T_0 + 1, \dots, T\}$. This means that calculating the duality gap at few random points would lead to the same type of guarantee with high probability. This approach has the advantage over averaging, since it is easier to implement the stopping condition (we simply check the duality gap at some random stopping points. This is in contrast to averaging in which we need to know T, T_0 in advance).

Remark 4 The above theorem applies to the hinge-loss function, $\phi_i(u) = \max\{0, 1 - y_i u\}$. However, for the hinge-loss, the constant 4 in the first inequality can be replaced by 1 (this is because the domain of the dual variables is positive, hence the constant 4 in Lemma 22 can be replaced by 1). We therefore obtain the bound:

$$T \geq T_0 + n + \frac{L^2}{\lambda\epsilon_P} \geq \max(0, \lceil n \log(0.5\lambda n L^{-2}) \rceil) + n + \frac{5L^2}{\lambda\epsilon_P}.$$

Theorem 5 Consider Procedure SDCA with $\alpha^{(0)} = 0$. Assume that ϕ_i is $(1/\gamma)$ -smooth for all i . To obtain an expected duality gap of $\mathbb{E}[P(w^{(T)}) - D(\alpha^{(T)})] \leq \epsilon_P$, it suffices to have a total number of iterations of

$$T \geq \left(n + \frac{1}{\lambda\gamma}\right) \log\left(\left(n + \frac{1}{\lambda\gamma}\right) \cdot \frac{1}{\epsilon_P}\right).$$

Moreover, to obtain an expected duality gap of $\mathbb{E}[P(\bar{w}) - D(\bar{\alpha})] \leq \epsilon_P$, it suffices to have a total number of iterations of $T > T_0$ where

$$T_0 \geq \left(n + \frac{1}{\lambda\gamma}\right) \log\left(\left(n + \frac{1}{\lambda\gamma}\right) \cdot \frac{1}{(T-T_0)\epsilon_P}\right).$$

Remark 6 If we choose $T = 2T_0$, and assume that $T_0 \geq n + 1/(\lambda\gamma)$, then the second part of Theorem 5 implies a requirement of

$$T_0 \geq \left(n + \frac{1}{\lambda\gamma}\right) \log\left(\frac{1}{\epsilon_P}\right),$$

which is slightly weaker than the first part of Theorem 5 when ϵ_P is relatively large.

Remark 7 *Bottou and Bousquet (2008) analyzed the runtime of SGD and other algorithms from the perspective of the time required to achieve a certain level of error on the test set. To perform such analysis, we also need to take into account the estimation error, namely, the additional error we suffer due to the fact that the training examples defining the regularized loss minimization problem are only a finite sample from the underlying distribution. The estimation error of the primal objective behaves like $\Theta(\frac{1}{\lambda n})$ (see Shalev-Shwartz and Srebro, 2008; Sridharan et al., 2009). Therefore, an interesting regime is when $\frac{1}{\lambda n} = \Theta(\epsilon)$. In that case, the bound for both Lipschitz and smooth functions would be $\tilde{O}(n)$. However, this bound on the estimation error is for the worst-case distribution over examples. Therefore, another interesting regime is when we would like $\epsilon \ll \frac{1}{\lambda n}$, but still $\frac{1}{\lambda n} = O(1)$ (following the practical observation that $\lambda = \Theta(1/n)$ often performs well). In that case, smooth functions still yield the bound $\tilde{O}(n)$, but the dominating term for Lipschitz functions will be $\frac{1}{\lambda \epsilon}$.*

Remark 8 *The runtime of SGD is $\tilde{O}(\frac{1}{\lambda \epsilon})$. This can be better than SDCA if $n \gg \frac{1}{\lambda \epsilon}$. However, in that case, SGD in fact only looks at $n' = \tilde{O}(\frac{1}{\lambda \epsilon})$ examples, so we can run SDCA on these n' examples and obtain basically the same rate. For smooth functions, SGD can be much worse than SDCA if $\epsilon \ll \frac{1}{\lambda n}$.*

4. Using SGD At The First Epoch

From the convergence analysis, SDCA may not perform as well as SGD for the first few epochs (each epoch means one pass over the data). The main reason is that SGD takes a larger step size than SDCA earlier on, which helps its performance. It is thus natural to combine SGD and SDCA, where the first epoch is performed using a modified stochastic gradient descent rule. We show that the expected dual sub-optimality at the end of the first epoch is $\tilde{O}(1/(\lambda n))$. This result can be combined with SDCA to obtain a faster convergence when $\lambda \gg \log n/n$.

We first introduce convenient notation. Let P_t denote the primal objective for the first t examples in the training set,

$$P_t(w) = \left[\frac{1}{t} \sum_{i=1}^t \phi_i(w^\top x_i) + \frac{\lambda}{2} \|w\|^2 \right].$$

The corresponding dual objective is

$$D_t(\alpha) = \left[\frac{1}{t} \sum_{i=1}^t -\phi_i^*(-\alpha_i) - \frac{\lambda}{2} \left\| \frac{1}{\lambda t} \sum_{i=1}^t \alpha_i x_i \right\|^2 \right].$$

Note that $P_n(w)$ is the primal objective given in (1) and that $D_n(\alpha)$ is the dual objective given in (2).

The following algorithm is a modification of SGD. The idea is to greedily decrease the dual sub-optimality for problem $D_t(\cdot)$ at each step t . This is different from DCA which works with $D_n(\cdot)$ at each step t .

Procedure Modified-SGD

Initialize: $w^{(0)} = 0$
Iterate: for $t = 1, 2, \dots, n$:
 Find α_t to maximize $-\phi_t^*(-\alpha_t) - \frac{\lambda t}{2} \|w^{(t-1)} + (\lambda t)^{-1} \alpha_t x_t\|^2$.
 Let $w^{(t)} = \frac{1}{\lambda t} \sum_{i=1}^t \alpha_i x_i$
 return α

We have the following result for the convergence of dual objective:

Theorem 9 Assume that ϕ_i is L -Lipschitz for all i . In addition, assume that (ϕ_i, x_i) are iid samples from the same distribution for all $i = 1, \dots, n$. At the end of Procedure Modified-SGD, we have

$$\mathbb{E}[D(\alpha^*) - D(\alpha)] \leq \frac{2L^2 \log(en)}{\lambda n}.$$

Here the expectation is with respect to the random sampling of $\{(\phi_i, x_i) : i = 1, \dots, n\}$.

Remark 10 When λ is relatively large, the convergence rate in Theorem 9 for modified-SGD is better than what we can prove for SDCA. This is because Modified-SGD employs a larger step size at each step t for $D_t(\alpha)$ than the corresponding step size in SDCA for $D(\alpha)$. However, the proof requires us to assume that (ϕ_i, x_i) are randomly drawn from a certain distribution, while this extra randomness assumption is not needed for the convergence of SDCA.

Procedure SDCA with SGD Initialization

Stage 1: call Procedure Modified-SGD and obtain α
Stage 2: call Procedure SDCA with parameter $\alpha^{(0)} = \alpha$

Theorem 11 Assume that ϕ_i is L -Lipschitz for all i . In addition, assume that (ϕ_i, x_i) are iid samples from the same distribution for all $i = 1, \dots, n$. Consider Procedure SDCA with SGD Initialization. To obtain a duality gap of $\mathbb{E}[P(\bar{w}) - D(\bar{\alpha})] \leq \epsilon_P$ at Stage 2, it suffices to have a total number of SDCA iterations of

$$T \geq T_0 + n + \frac{4L^2}{\lambda \epsilon_P} \geq \lceil n \log(\log(en)) \rceil + n + \frac{20L^2}{\lambda \epsilon_P}.$$

Moreover, when $t \geq T_0$, we have duality sub-optimality bound of $\mathbb{E}[D(\alpha^*) - D(\alpha^{(t)})] \leq \epsilon_P/2$.

Remark 12 For Lipschitz loss, ideally we would like to have a computational complexity of $O(n + L^2/(\lambda \epsilon_P))$. Theorem 11 shows that SDCA with SGD at first epoch can achieve no worst than $O(n \log(\log n) + L^2/(\lambda \epsilon_P))$, which is very close to the ideal bound. The result is better than that of vanilla SDCA in Theorem 2 when λ is relatively large, which shows a complexity of $O(n \log(n) + L^2/(\lambda \epsilon_P))$. The difference is caused by small step-sizes in the vanilla SDCA, and its negative effect can be observed in practice. That is, the vanilla SDCA tends to have a slower convergence rate than SGD in the first few iterations when λ is relatively large.

Remark 13 *Similar to Remark 4, for the hinge-loss, the constant 4 in Theorem 11 can be reduced to 1, and the constant 20 can be reduced to 5.*

5. Refined Analysis For Almost Smooth Loss

Our analysis shows that for smooth loss, SDCA converges faster than SGD (linear versus sub-linear convergence). For non-smooth loss, the analysis does not show any advantage of SDCA over SGD. This does not explain the practical observation that SDCA converges faster than SGD asymptotically even for SVM. This section tries to refine the analysis for Lipschitz loss and shows potential advantage of SDCA over SGD asymptotically. Note that the refined analysis of this section relies on quantities that depend on the underlying data distribution, and thus the results are more complicated than those presented earlier. Although precise interpretations of these results will be complex, we will discuss them qualitatively after the theorem statements, and use them to explain the advantage of SDCA over SGD for non-smooth losses.

Although we note that for SVM, Luo and Tseng's analysis (Luo and Tseng, 1992) shows linear convergence of the form $(1 - \nu)^k$ for dual sub-optimality after k passes over the data, as we mentioned, ν is proportional to the smallest nonzero eigenvalue of the data Gram matrix $X^\top X$, and hence can be arbitrarily bad when two data points $x_i \neq x_j$ becomes very close to each other. Our analysis uses a completely different argument that avoids this dependency on the data Gram matrix.

The main intuition behind our analysis is that many non-smooth loss functions are nearly smooth everywhere. For example, the hinge loss $\max(0, 1 - uy_i)$ is smooth at any point u such that uy_i is not close to 1. Since a smooth loss has a strongly convex dual (and the strong convexity of the dual is directly used in our proof to obtain fast rate for smooth loss), the refined analysis in this section relies on the following refined dual strong convexity condition that holds for nearly everywhere smooth loss functions.

Definition 14 *For each i , we define $\gamma_i(\cdot) \geq 0$ so that for all dual variables a and b , and $u \in \partial\phi_i^*(-b)$, we have*

$$\phi_i^*(-a) - \phi_i^*(-b) + u(a - b) \geq \gamma_i(u)|a - b|^2. \quad (4)$$

For the SVM loss, we have $\phi_i(u) = \max(0, 1 - uy_i)$, and $\phi_i^*(-a) = -ay_i$, with $ay_i \in [0, 1]$ and $y_i \in \{\pm 1\}$. It follows that

$$\phi_i^*(-a) - \phi_i^*(-b) + u(a - b) = (b - a)y_i + u(a - b) = |uy_i - 1||a - b| \geq |uy_i - 1| \cdot |a - b|^2.$$

Therefore we may take $\gamma_i(u) = |uy_i - 1|$.

For the absolute deviation loss, we have $\phi_i(u) = |u - y_i|$, and $\phi^*(-a) = -ay_i$ with $a \in [-1, 1]$. It follows that $\gamma_i(u) = |u - y_i|$.

Proposition 15 *Under the assumption of (4). Let $\gamma_i = \gamma_i(w^{*\top} x_i)$, we have the following dual strong convexity inequality:*

$$D(\alpha^*) - D(\alpha) \geq \frac{1}{n} \sum_{i=1}^n \gamma_i |\alpha_i - \alpha_i^*|^2 + \frac{\lambda}{2} (w - w^*)^\top (w - w^*). \quad (5)$$

Moreover, given $w \in \mathbb{R}^d$ and $-a_i \in \partial\phi_i(w^\top x_i)$, we have

$$|(w^* - w)^\top x_i| \geq \gamma_i |a_i - \alpha_i^*|.$$

For SVM, we can take $\gamma_i = |w^{*\top} x_i y_i - 1|$, and for the absolute deviation loss, we may take $\gamma_i = |w^{*\top} x_i - y_i|$. Although some of γ_i can be close to zero, in practice, most γ_i will be away from zero, which means $D(\alpha)$ is strongly convex at nearly all points. Under this assumption, we may establish a convergence result for the dual sub-optimality.

Theorem 16 *Consider Procedure SDCA with $\alpha^{(0)} = 0$. Assume that ϕ_i is L -Lipschitz for all i and it satisfies (5). Define $N(u) = \#\{i : \gamma_i < u\}$. To obtain a dual-suboptimality of $\mathbb{E}[D(\alpha^*) - D(\alpha^t)] \leq \epsilon_D$, it suffices to have a total number of iterations of*

$$t \geq 2(n/s) \log(2/\epsilon_D),$$

where $s \in [0, 1]$ satisfies $\epsilon_D \geq 8L^2(s/\lambda n)N(s/\lambda n)/n$.

Remark 17 *if $N(s/\lambda n)/n$ is small, then Theorem 16 is superior to Theorem 2 for the convergence of the dual objective function. We consider three scenarios. The first scenario is when $s = 1$. If $N(1/\lambda n)/n$ is small, and $\epsilon_D \geq 8L^2(1/\lambda n)N(1/\lambda n)/n$, then the convergence is linear. The second scenario is when there exists s_0 so that $N(s_0/\lambda n) = 0$ (for SVM, it means that $\lambda n|w^{*\top} x_i y_i - 1| \geq s_0$ for all i), and since $\epsilon_D \geq 8L^2(s_0/\lambda n)N(s_0/\lambda n)/n$, we again have a linear convergence of $(2n/s_0) \log(2/\epsilon_D)$. In the third scenario, we assume that $N(s/\lambda n)/n = O[(s/\lambda n)^\nu]$ for some $\nu > 0$, we can take $\epsilon_D = O((s/\lambda n)^{1+\nu})$ and obtain*

$$t \geq O(\lambda^{-1} \epsilon_D^{-1/(1+\nu)} \log(2/\epsilon_D)).$$

The $\log(1/\epsilon_D)$ factor can be removed in this case with a slightly more complex analysis. This result is again superior to Theorem 2 for dual convergence.

The following result shows fast convergence of duality gap using Theorem 16.

Theorem 18 *Consider Procedure SDCA with $\alpha^{(0)} = 0$. Assume that ϕ_i is L -Lipschitz for all i and it satisfies (4). Let $\rho \leq 1$ be the largest eigenvalue of the matrix $n^{-1} \sum_{i=1}^n x_i x_i^\top$. Define $N(u) = \#\{i : \gamma_i < u\}$. Assume that at time $T_0 \geq n$, we have dual suboptimality of $\mathbb{E}[D(\alpha^*) - D(\alpha^{(T_0)})] \leq \epsilon_D$, and define*

$$\tilde{\epsilon}_P = \inf_{\gamma > 0} \left[\frac{N(\gamma)}{n} 4L^2 + \frac{2\epsilon_D}{\min(\gamma, \lambda\gamma^2/(2\rho))} \right],$$

then at time $T = 2T_0$, we have

$$\mathbb{E}[P(\bar{w}) - D(\bar{\alpha})] \leq \epsilon_D + \frac{\tilde{\epsilon}_P}{2\lambda T_0}.$$

If for some γ , $N(\gamma)/n$ is small, then Theorem 18 is superior to Theorem 2. Although the general dependency may be complex, the improvement over Theorem 2 can be more easily seen in the special case that $N(\gamma) = 0$ for some $\gamma > 0$. In fact, in this case we have $\tilde{\epsilon}_P = O(\epsilon_D)$, and thus

$$\mathbb{E}[P(\bar{w}) - D(\bar{\alpha})] = O(\epsilon_D).$$

This means that the convergence rate for duality gap in Theorem 18 is linear as implied by the linear convergence of ϵ_D in Theorem 16.

6. Examples

We will specify the SDCA algorithms for a few common loss functions. For simplicity, we only specify the algorithms without SGD initialization. In practice, instead of complete randomization, we may also run in epochs, and each epoch employs a random permutation of the data. We call this variant SDCA-Perm.

Procedure SDCA-Perm($\alpha^{(0)}$)

Let $w^{(0)} = w(\alpha^{(0)})$
Let $t = 0$
Iterate: for epoch $k = 1, 2, \dots$
 Let $\{i_1, \dots, i_n\}$ be a random permutation of $\{1, \dots, n\}$
 Iterate: for $j = 1, 2, \dots, n$:
 $t \leftarrow t + 1$
 $i = i_j$
 Find $\Delta\alpha_i$ to increase dual (*)
 $\alpha^{(t)} \leftarrow \alpha^{(t-1)} + \Delta\alpha_i e_i$
 $w^{(t)} \leftarrow w^{(t-1)} + (\lambda n)^{-1} \Delta\alpha_i x_i$
Output (Averaging option):
 Let $\bar{\alpha} = \frac{1}{T-T_0} \sum_{i=T_0+1}^T \alpha^{(i-1)}$
 Let $\bar{w} = w(\bar{\alpha}) = \frac{1}{T-T_0} \sum_{i=T_0+1}^T w^{(i-1)}$
 return \bar{w}
Output (Random option):
 Let $\bar{\alpha} = \alpha^{(t)}$ and $\bar{w} = w^{(t)}$ for some random $t \in T_0 + 1, \dots, T$
 return \bar{w}

6.1 Lipschitz Loss

Hinge loss is used in SVM. We have $\phi_i(u) = \max\{0, 1 - y_i u\}$ and $\phi_i^*(-a) = -ay_i$ with $ay_i \in [0, 1]$. Absolute deviation loss is used in quantile regression. We have $\phi_i(u) = |u - y_i|$ and $\phi_i^*(-a) = -ay_i$ with $a \in [-1, 1]$.

For the hinge loss, step (*) in Procedure SDCA-Perm has a closed form solution as

$$\Delta\alpha_i = y_i \max \left(0, \min \left(1, \frac{1 - x_i^\top w^{(t-1)} y_i}{\|x_i\|^2 / (\lambda n)} + \alpha_i^{(t-1)} y_i \right) \right) - \alpha_i^{(t-1)}.$$

For absolute deviation loss, step (*) in Procedure SDCA-Perm has a closed form solution as

$$\Delta\alpha_i = \max \left(-1, \min \left(1, \frac{y_i - x_i^\top w^{(t-1)}}{\|x_i\|^2 / (\lambda n)} + \alpha_i^{(t-1)} \right) \right) - \alpha_i^{(t-1)}.$$

Both hinge loss and absolute deviation loss are 1-Lipschitz. Therefore, we expect a convergence behavior of no worse than

$$O \left(n \log n + \frac{1}{\lambda \epsilon} \right)$$

without SGD initialization based on Theorem 2. The refined analysis in Section 5 suggests a rate that can be significantly better, and this is confirmed with our empirical experiments.

6.2 Smooth Loss

Squared loss is used in ridge regression. We have $\phi_i(u) = (u - y_i)^2$, and $\phi_i^*(-a) = -ay_i + a^2/4$. Log loss is used in logistic regression. We have $\phi_i(u) = \log(1 + \exp(-y_i u))$, and $\phi_i^*(-a) = ay_i \log(ay_i) + (1 - ay_i) \log(1 - ay_i)$ with $ay_i \in [0, 1]$.

For squared loss, step (*) in Procedure SDCA-Perm has a closed form solution as

$$\Delta\alpha_i = \frac{y_i - x_i^\top w^{(t-1)} - 0.5\alpha_i^{(t-1)}}{0.5 + \|x_i\|^2/(\lambda n)}.$$

For log loss, step (*) in Procedure SDCA-Perm does not have a closed form solution. However, one may start with the approximate solution,

$$\Delta\alpha_i = \frac{(1 + \exp(x_i^\top w^{(t-1)} y_i))^{-1} y_i - \alpha_i^{(t-1)}}{\max(1, 0.25 + \|x_i\|^2/(\lambda n))},$$

and further use several steps of Newton's update to get a more accurate solution.

Finally, we present a smooth variant of the hinge-loss, as defined below. Recall that the hinge loss function (for positive labels) is $\phi(u) = \max\{0, 1 - u\}$ and we have $\phi^*(-a) = -a$ with $a \in [0, 1]$. Consider adding to ϕ^* the term $\frac{\gamma}{2}a^2$ which yields the γ -strongly convex function

$$\tilde{\phi}_\gamma^*(a) = \phi^*(a) + \frac{\gamma}{2}a^2.$$

Then, its conjugate, which is defined below, is $(1/\gamma)$ -smooth. We refer to it as the *smoothed hinge-loss* (for positive labels):

$$\tilde{\phi}_\gamma(x) = \max_{a \in [-1, 0]} \left[ax - a - \frac{\gamma}{2}a^2 \right] = \begin{cases} 0 & x > 1 \\ 1 - x - \gamma/2 & x < 1 - \gamma \\ \frac{1}{2\gamma}(1 - x)^2 & \text{otherwise} \end{cases}. \quad (6)$$

For the smoothed hinge loss, step (*) in Procedure SDCA-Perm has a closed form solution as

$$\Delta\alpha_i = y_i \max \left(0, \min \left(1, \frac{1 - x_i^\top w^{(t-1)} y_i - \gamma \alpha_i^{(t-1)} y_i}{\|x_i\|^2/(\lambda n) + \gamma} + \alpha_i^{(t-1)} y_i \right) \right) - \alpha_i^{(t-1)}.$$

Both log loss and squared loss are 1-smooth. The smoothed-hinge loss is $1/\gamma$ smooth. Therefore we expect a convergence behavior of no worse than

$$O \left(\left(n + \frac{1}{\gamma \lambda} \right) \log \frac{1}{\epsilon} \right).$$

This is confirmed in our empirical experiments.

7. Proofs

We denote by $\partial\phi_i(a)$ the set of sub-gradients of ϕ_i at a . We use the notation $\phi'_i(a)$ to denote some sub-gradient of ϕ_i at a . For convenience, we list the following simple facts about primal and dual formulations, which will be used in the proofs. For each i , we have

$$-\alpha_i^* \in \partial\phi_i(w^{*\top} x_i), \quad w^{*\top} x_i \in \partial\phi_i^*(-\alpha_i^*),$$

and

$$w^* = \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i^* x_i.$$

The proof of our basic results stated in Theorem 5 and Theorem 2 relies on the fact that for SDCA, it is possible to lower bound the expected increase in dual objective by the duality gap. This key observation is stated in Lemma 19. Note that the duality gap can be further lower bounded using dual suboptimality. Therefore Lemma 19 implies a recursion for dual suboptimality which can be solved to obtain the convergence of dual objective. We can then apply Lemma 19 again, and the convergence of dual objective implies an upper bound of the duality gap, which leads to the basic theorems. The more refined results in Section 4 and Section 5 use similar strategies but with Lemma 19 replaced by its variants.

7.1 Proof Of Theorem 5

The key lemma, which estimates the expected increase in dual objective in terms of the duality gap, can be stated as follows.

Lemma 19 *Assume that ϕ_i^* is γ -strongly-convex (where γ can be zero). Then, for any iteration t and any $s \in [0, 1]$ we have*

$$\mathbb{E}[D(\alpha^{(t)}) - D(\alpha^{(t-1)})] \geq \frac{s}{n} \mathbb{E}[P(w^{(t-1)}) - D(\alpha^{(t-1)})] - \left(\frac{s}{n}\right)^2 \frac{G^{(t)}}{2\lambda},$$

where

$$G^{(t)} = \frac{1}{n} \sum_{i=1}^n \left(\|x_i\|^2 - \frac{\gamma(1-s)\lambda n}{s} \right) \mathbb{E}[(u_i^{(t-1)} - \alpha_i^{(t-1)})^2],$$

and $-u_i^{(t-1)} \in \partial\phi_i(x_i^\top w^{(t-1)})$.

Proof Since only the i 'th element of α is updated, the improvement in the dual objective can be written as

$$n[D(\alpha^{(t)}) - D(\alpha^{(t-1)})] = \underbrace{\left(-\phi_i^*(-\alpha_i^{(t)}) - \frac{\lambda n}{2} \|w^{(t)}\|^2 \right)}_A - \underbrace{\left(-\phi_i^*(-\alpha_i^{(t-1)}) - \frac{\lambda n}{2} \|w^{(t-1)}\|^2 \right)}_B.$$

By the definition of the update we have for all $s \in [0, 1]$ that

$$\begin{aligned} A &= \max_{\Delta\alpha_i} -\phi_i^*(-(\alpha_i^{(t-1)} + \Delta\alpha_i)) - \frac{\lambda n}{2} \|w^{(t-1)} + (\lambda n)^{-1} \Delta\alpha_i x_i\|^2 \\ &\geq -\phi_i^*(-(\alpha_i^{(t-1)} + s(u_i^{(t-1)} - \alpha_i^{(t-1)}))) - \frac{\lambda n}{2} \|w^{(t-1)} + (\lambda n)^{-1} s(u_i^{(t-1)} - \alpha_i^{(t-1)}) x_i\|^2. \quad (7) \end{aligned}$$

From now on, we omit the superscripts and subscripts. Since ϕ^* is γ -strongly convex, we have that

$$\phi^*(-(\alpha + s(u - \alpha))) = \phi^*(s(-u) + (1-s)(-\alpha)) \leq s\phi^*(-u) + (1-s)\phi^*(-\alpha) - \frac{\gamma}{2}s(1-s)(u - \alpha)^2.$$

Combining this with (7) and rearranging terms we obtain that

$$\begin{aligned} A &\geq -s\phi^*(-u) - (1-s)\phi^*(-\alpha) + \frac{\gamma}{2}s(1-s)(u - \alpha)^2 - \frac{\lambda n}{2}\|w + (\lambda n)^{-1}s(u - \alpha)x\|^2 \\ &= -s\phi^*(-u) - (1-s)\phi^*(-\alpha) + \frac{\gamma}{2}s(1-s)(u - \alpha)^2 - \frac{\lambda n}{2}\|w\|^2 - s(u - \alpha)w^\top x \\ &\quad - \frac{s^2(u - \alpha)^2}{2\lambda n}\|x\|^2 \\ &= \underbrace{-s(\phi^*(-u) + uw^\top x)}_{s\phi(w^\top x)} + \underbrace{(-\phi^*(-\alpha) - \frac{\lambda n}{2}\|w\|^2)}_B + \frac{s}{2}\left(\gamma(1-s) - \frac{s\|x\|^2}{\lambda n}\right)(u - \alpha)^2 \\ &\quad + s(\phi^*(-\alpha) + \alpha w^\top x), \end{aligned}$$

where we used $-u \in \partial\phi(w^\top x)$ which yields $\phi^*(-u) = -uw^\top x - \phi(w^\top x)$. Therefore

$$A - B \geq s \left[\phi(w^\top x) + \phi^*(-\alpha) + \alpha w^\top x + \left(\frac{\gamma(1-s)}{2} - \frac{s\|x\|^2}{2\lambda n} \right) (u - \alpha)^2 \right]. \quad (8)$$

Next note that

$$\begin{aligned} P(w) - D(\alpha) &= \frac{1}{n} \sum_{i=1}^n \phi_i(w^\top x_i) + \frac{\lambda}{2} w^\top w - \left(-\frac{1}{n} \sum_{i=1}^n \phi_i^*(-\alpha_i) - \frac{\lambda}{2} w^\top w \right) \\ &= \frac{1}{n} \sum_{i=1}^n \left(\phi_i(w^\top x_i) + \phi_i^*(-\alpha_i) + \alpha_i w^\top x_i \right). \end{aligned}$$

Therefore, if we take expectation of (8) w.r.t. the choice of i we obtain that

$$\frac{1}{s} \mathbb{E}[A - B] \geq \mathbb{E}[P(w) - D(\alpha)] - \frac{s}{2\lambda n} \cdot \underbrace{\frac{1}{n} \sum_{i=1}^n \left(\|x_i\|^2 - \frac{\gamma(1-s)\lambda n}{s} \right) \mathbb{E}(u_i - \alpha_i)^2}_{=G^{(t)}}.$$

We have obtained that

$$\frac{n}{s} \mathbb{E}[D(\alpha^{(t)}) - D(\alpha^{(t-1)})] \geq \mathbb{E}[P(w^{(t-1)}) - D(\alpha^{(t-1)})] - \frac{sG^{(t)}}{2\lambda n}.$$

Multiplying both sides by s/n concludes the proof of the lemma. ■

We also use the following simple lemma:

Lemma 20 *For all α , $D(\alpha) \leq P(w^*) \leq P(0) \leq 1$. In addition, $D(0) \geq 0$.*

Proof The first inequality is by weak duality, the second is by the optimality of w^* , and the third by the assumption that $\phi_i(0) \leq 1$. For the last inequality we use $-\phi_i^*(0) = -\max_z(0 - \phi_i(z)) = \min_z \phi_i(z) \geq 0$, which yields $D(0) \geq 0$. ■

Equipped with the above lemmas we are ready to prove Theorem 5.

Proof [Proof of Theorem 5] The assumption that ϕ_i is $(1/\gamma)$ -smooth implies that ϕ_i^* is γ -strongly-convex. We will apply Lemma 19 with $s = \frac{\lambda n \gamma}{1 + \lambda n \gamma} \in [0, 1]$. Recall that $\|x_i\| \leq 1$. Therefore, the choice of s implies that $\|x_i\|^2 - \frac{\gamma(1-s)\lambda n}{s} \leq 0$, and hence $G^{(t)} \leq 0$ for all t . This yields,

$$\mathbb{E}[D(\alpha^{(t)}) - D(\alpha^{(t-1)})] \geq \frac{s}{n} \mathbb{E}[P(w^{(t-1)}) - D(\alpha^{(t-1)})] .$$

But since $\epsilon_D^{(t-1)} := D(\alpha^*) - D(\alpha^{(t-1)}) \leq P(w^{(t-1)}) - D(\alpha^{(t-1)})$ and $D(\alpha^{(t)}) - D(\alpha^{(t-1)}) = \epsilon_D^{(t-1)} - \epsilon_D^{(t)}$, we obtain that

$$\mathbb{E}[\epsilon_D^{(t)}] \leq (1 - \frac{s}{n}) \mathbb{E}[\epsilon_D^{(t-1)}] \leq (1 - \frac{s}{n})^t \mathbb{E}[\epsilon_D^{(0)}] \leq (1 - \frac{s}{n})^t \leq \exp(-st/n) = \exp\left(-\frac{\lambda \gamma t}{1 + \lambda \gamma n}\right) .$$

This would be smaller than ϵ_D if

$$t \geq \left(n + \frac{1}{\lambda \gamma}\right) \log(1/\epsilon_D) .$$

It implies that

$$\mathbb{E}[P(w^{(t)}) - D(\alpha^{(t)})] \leq \frac{n}{s} \mathbb{E}[\epsilon_D^{(t)} - \epsilon_D^{(t+1)}] \leq \frac{n}{s} \mathbb{E}[\epsilon_D^{(t)}]. \quad (9)$$

So, requiring $\epsilon_D^{(t)} \leq \frac{s}{n} \epsilon_P$ we obtain a duality gap of at most ϵ_P . This means that we should require

$$t \geq \left(n + \frac{1}{\lambda \gamma}\right) \log\left((n + \frac{1}{\lambda \gamma}) \cdot \frac{1}{\epsilon_P}\right) ,$$

which proves the first part of Theorem 5.

Next, we sum (9) over $t = T_0, \dots, T-1$ to obtain

$$\mathbb{E}\left[\frac{1}{T-T_0} \sum_{t=T_0}^{T-1} (P(w^{(t)}) - D(\alpha^{(t)}))\right] \leq \frac{n}{s(T-T_0)} \mathbb{E}[D(\alpha^{(T)}) - D(\alpha^{(T_0)})].$$

Now, if we choose $\bar{w}, \bar{\alpha}$ to be either the average vectors or a randomly chosen vector over $t \in \{T_0 + 1, \dots, T\}$, then the above implies

$$\mathbb{E}[P(\bar{w}) - D(\bar{\alpha})] \leq \frac{n}{s(T-T_0)} \mathbb{E}[D(\alpha^{(T)}) - D(\alpha^{(T_0)})] \leq \frac{n}{s(T-T_0)} \mathbb{E}[\epsilon_D^{(T_0)}].$$

It follows that in order to obtain a result of $\mathbb{E}[P(\bar{w}) - D(\bar{\alpha})] \leq \epsilon_P$, we only need to have

$$\mathbb{E}[\epsilon_D^{(T_0)}] \leq \frac{s(T-T_0)\epsilon_P}{n} = \frac{(T-T_0)\epsilon_P}{n + \frac{1}{\lambda \gamma}}.$$

This implies the second part of Theorem 5, and concludes the proof. ■

7.2 Proof Of Theorem 2

Next, we turn to the case of Lipschitz loss function. We rely on the following lemma.

Lemma 21 *Let $\phi : \mathbb{R} \rightarrow \mathbb{R}$ be an L -Lipschitz function. Then, for any α s.t. $|\alpha| > L$ we have that $\phi^*(\alpha) = \infty$.*

Proof Fix some $\alpha > L$. By definition of the conjugate we have

$$\begin{aligned} \phi^*(\alpha) &= \sup_x [\alpha x - \phi(x)] \\ &\geq -\phi(0) + \sup_x [\alpha x - (\phi(x) - \phi(0))] \\ &\geq -\phi(0) + \sup_x [\alpha x - L|x - 0|] \\ &\geq -\phi(0) + \sup_{x>0} (\alpha - L)x = \infty. \end{aligned}$$

Similar argument holds for $\alpha < -L$. ■

A direct corollary of the above lemma is:

Lemma 22 *Suppose that for all i , ϕ_i is L -Lipschitz. Let $G^{(t)}$ be as defined in Lemma 19 (with $\gamma = 0$). Then, $G^{(t)} \leq 4L^2$.*

Proof Using Lemma 21 we know that $|\alpha_i^{(t-1)}| \leq L$, and in addition by the relation of Lipschitz and sub-gradients we have $|u_i^{(t-1)}| \leq L$. Thus, $(u_i^{(t-1)} - \alpha_i^{(t-1)})^2 \leq 4L^2$, and the proof follows. ■

We are now ready to prove Theorem 2.

Proof [Proof of Theorem 2] Let $G = \max_t G^{(t)}$ and note that by Lemma 22 we have $G \leq 4L^2$. Lemma 19, with $\gamma = 0$, tells us that

$$\mathbb{E}[D(\alpha^{(t)}) - D(\alpha^{(t-1)})] \geq \frac{s}{n} \mathbb{E}[P(w^{(t-1)}) - D(\alpha^{(t-1)})] - \left(\frac{s}{n}\right)^2 \frac{G}{2\lambda}, \quad (10)$$

which implies that

$$\mathbb{E}[\epsilon_D^{(t)}] \leq \left(1 - \frac{s}{n}\right) \mathbb{E}[\epsilon_D^{(t-1)}] + \left(\frac{s}{n}\right)^2 \frac{G}{2\lambda}.$$

We next show that the above yields

$$\mathbb{E}[\epsilon_D^{(t)}] \leq \frac{2G}{\lambda(2n + t - t_0)} \quad (11)$$

for all $t \geq t_0 = \max(0, \lceil n \log(2\lambda n \epsilon_D^{(0)} / G) \rceil)$. Indeed, let us choose $s = 1$, then at $t = t_0$, we have

$$\mathbb{E}[\epsilon_D^{(t)}] \leq \left(1 - \frac{1}{n}\right)^t \epsilon_D^{(0)} + \frac{G}{2\lambda n^2} \frac{1}{1 - (1 - 1/n)} \leq e^{-t/n} \epsilon_D^{(0)} + \frac{G}{2\lambda n} \leq \frac{G}{\lambda n}.$$

This implies that (11) holds at $t = t_0$. For $t > t_0$ we use an inductive argument. Suppose the claim holds for $t - 1$, therefore

$$\mathbb{E}[\epsilon_D^{(t)}] \leq \left(1 - \frac{s}{n}\right) \mathbb{E}[\epsilon_D^{(t-1)}] + \left(\frac{s}{n}\right)^2 \frac{G}{2\lambda} \leq \left(1 - \frac{s}{n}\right) \frac{2G}{\lambda(2n + t - 1 - t_0)} + \left(\frac{s}{n}\right)^2 \frac{G}{2\lambda}.$$

Choosing $s = 2n/(2n - t_0 + t - 1) \in [0, 1]$ yields

$$\begin{aligned}
 \mathbb{E}[\epsilon_D^{(t)}] &\leq \left(1 - \frac{2}{2n-t_0+t-1}\right) \frac{2G}{\lambda(2n-t_0+t-1)} + \left(\frac{2}{2n-t_0+t-1}\right)^2 \frac{G}{2\lambda} \\
 &= \frac{2G}{\lambda(2n-t_0+t-1)} \left(1 - \frac{1}{2n-t_0+t-1}\right) \\
 &= \frac{2G}{\lambda(2n-t_0+t-1)} \frac{2n-t_0+t-2}{2n-t_0+t-1} \\
 &\leq \frac{2G}{\lambda(2n-t_0+t-1)} \frac{2n-t_0+t-1}{2n-t_0+t} \\
 &= \frac{2G}{\lambda(2n-t_0+t)}.
 \end{aligned}$$

This provides a bound on the dual sub-optimality. We next turn to bound the duality gap. Summing (10) over $t = T_0 + 1, \dots, T$ and rearranging terms we obtain that

$$\mathbb{E} \left[\frac{1}{T - T_0} \sum_{t=T_0+1}^T (P(w^{(t-1)}) - D(\alpha^{(t-1)})) \right] \leq \frac{n}{s(T - T_0)} \mathbb{E}[D(\alpha^{(T)}) - D(\alpha^{(T_0)})] + \frac{sG}{2\lambda n}.$$

Now, if we choose $\bar{w}, \bar{\alpha}$ to be either the average vectors or a randomly chosen vector over $t \in \{T_0 + 1, \dots, T\}$, then the above implies

$$\mathbb{E}[P(\bar{w}) - D(\bar{\alpha})] \leq \frac{n}{s(T - T_0)} \mathbb{E}[D(\alpha^{(T)}) - D(\alpha^{(T_0)})] + \frac{sG}{2\lambda n}.$$

If $T \geq n + T_0$ and $T_0 \geq t_0$, we can set $s = n/(T - T_0)$ and combining with (11) we obtain

$$\begin{aligned}
 \mathbb{E}[P(\bar{w}) - D(\bar{\alpha})] &\leq \mathbb{E}[D(\alpha^{(T)}) - D(\alpha^{(T_0)})] + \frac{G}{2\lambda(T - T_0)} \\
 &\leq \mathbb{E}[D(\alpha^*) - D(\alpha^{(T_0)})] + \frac{G}{2\lambda(T - T_0)} \\
 &\leq \frac{2G}{\lambda(2n - t_0 + T_0)} + \frac{G}{2\lambda(T - T_0)}.
 \end{aligned}$$

A sufficient condition for the above to be smaller than ϵ_P is that $T_0 \geq \frac{4G}{\lambda\epsilon_P} - 2n + t_0$ and $T \geq T_0 + \frac{G}{\lambda\epsilon_P}$. It also implies that $\mathbb{E}[D(\alpha^*) - D(\alpha^{(T_0)})] \leq \epsilon_P/2$. Since we also need $T_0 \geq t_0$ and $T - T_0 \geq n$, the overall number of required iterations can be

$$T_0 \geq \max\{t_0, 4G/(\lambda\epsilon_P) - 2n + t_0\}, \quad T - T_0 \geq \max\{n, G/(\lambda\epsilon_P)\}.$$

We conclude the proof by noticing that $\epsilon_D^{(0)} \leq 1$ using Lemma 20, which implies that $t_0 \leq \max(0, \lceil n \log(2\lambda n/G) \rceil)$. ■

7.3 Proof Of Theorem 9

We assume that (ϕ_t, x_t) are randomly drawn from a distribution D , and define the population optimizer

$$w_D^* = \operatorname{argmin}_w P_D(w), \quad P_D(w) = \mathbb{E}_{(\phi, x) \sim D} \left[\phi(w^\top x) + \frac{\lambda}{2} \|w\|^2 \right].$$

By definition, we have $P(w^*) \leq P(w_D^*)$ for any specific realization of $\{(\phi_t, x_t) : t = 1, \dots, n\}$. Therefore

$$\mathbb{E} P(w^*) \leq \mathbb{E} P(w_D^*) = \mathbb{E} P_D(w_D^*),$$

where the expectation is with respect to the choice of examples, and note that both $P(\cdot)$ and w^* are sample dependent.

After each step t , we let $\alpha^{(t)} = [\alpha_1, \dots, \alpha_t]$, and let $-u \in \partial \phi_{t+1}(x_{t+1}^\top w^{(t)})$. We have, for all t ,

$$\begin{aligned} (t+1)D_{t+1}(\alpha^{(t+1)}) - tD_t(\alpha^{(t)}) &= -\phi_{t+1}^*(-\alpha_{t+1}^{(t+1)}) - (t+1)\frac{\lambda}{2}\|w^{(t+1)}\|^2 + t\frac{\lambda}{2}\|w^{(t)}\|^2 \\ &= -\phi_{t+1}^*(-\alpha_{t+1}^{(t+1)}) - \frac{1}{2(t+1)\lambda}\|\lambda t w^{(t)} + \alpha_{t+1}^{(t+1)} x_{t+1}\|^2 + \frac{1}{2t\lambda}\|\lambda t w^{(t)}\|^2 \\ &\geq -\phi_{t+1}^*(-u) - \frac{1}{2(t+1)\lambda}\|\lambda t w^{(t)} + u x_{t+1}\|^2 + \frac{1}{2t\lambda}\|\lambda t w^{(t)}\|^2 \\ &= -\phi_{t+1}^*(-u) - \frac{t}{t+1} x_{t+1}^\top w^{(t)} u + \frac{1}{2\lambda} \left(\frac{1}{t} - \frac{1}{t+1} \right) \|\lambda t w^{(t)}\|^2 - \frac{u^2 \|x_{t+1}\|^2}{2(t+1)\lambda} \\ &= -\phi_{t+1}^*(-u) - x_{t+1}^\top w^{(t)} u + \left(1 - \frac{t}{t+1} \right) x_{t+1}^\top w^{(t)} u + \frac{1}{2(t+1)\lambda} \left(\frac{\|\lambda t w^{(t)}\|^2}{t} - u^2 \|x_{t+1}\|^2 \right) \\ &= \phi_{t+1}(x_{t+1}^\top w^{(t)}) + \frac{1}{2(t+1)\lambda} \left(2\lambda x_{t+1}^\top w^{(t)} u + \frac{\|\lambda t w^{(t)}\|^2}{t} - u^2 \|x_{t+1}\|^2 \right) \\ &= \phi_{t+1}(x_{t+1}^\top w^{(t)}) + \frac{\lambda}{2} \|w^{(t)}\|^2 + \frac{1}{2(t+1)\lambda} \left(2\lambda x_{t+1}^\top w^{(t)} u - \|\lambda w^{(t)}\|^2 - u^2 \|x_{t+1}\|^2 \right) \\ &= \phi_{t+1}(w^{(t) \top} x_{t+1}) + \frac{\lambda}{2} \|w^{(t)}\|^2 - \frac{\|\lambda w^{(t)} - u x_{t+1}\|^2}{2(t+1)\lambda}. \end{aligned}$$

The inequality above can be obtained by noticing that the choice of $-\alpha_{t+1}^{(t+1)}$ maximizes the dual objective. In the derivation of the equalities we have used basic algebra as well as the equation $-\phi_{t+1}^*(-u) - x_{t+1}^\top w^{(t)} u = \phi_{t+1}(x_{t+1}^\top w^{(t)})$ which follows from $-u \in \partial \phi_{t+1}(x_{t+1}^\top w^{(t)})$. Next we note that $\|\lambda w^{(t)} - u x_{t+1}\| \leq 2L$ (where we used the triangle inequality, the definition of $w^{(t)}$, and Lemma 21). Therefore,

$$(t+1)D_{t+1}(\alpha^{(t+1)}) - tD_t(\alpha^{(t)}) \geq \phi_{t+1}(w^{(t) \top} x_{t+1}) + \frac{\lambda}{2} \|w^{(t)}\|^2 - \frac{2L^2}{(t+1)\lambda}.$$

Taking expectation *with respect to the choice of the examples*, and note that the $(t+1)$ 'th example does not depend on $w^{(t)}$ we obtain that

$$\begin{aligned} &\mathbb{E}[(t+1)D_{t+1}(\alpha^{(t+1)}) - tD_t(\alpha^{(t)})] \\ &\geq \mathbb{E}[P_D(w^{(t)})] - \frac{2L^2}{(t+1)\lambda} \geq \mathbb{E}[P_D(w_D^*)] - \frac{2L^2}{(t+1)\lambda} \\ &\geq \mathbb{E}[P(w^*)] - \frac{2L^2}{(t+1)\lambda} = \mathbb{E}[D(\alpha^*)] - \frac{2L^2}{(t+1)\lambda}. \end{aligned}$$

Using Lemma 20 we know that $D_t(\alpha^{(t)}) \geq 0$ for all t . Therefore, by summing the above over t we obtain that

$$\mathbb{E}[nD(\alpha^{(n)})] \geq n \mathbb{E}[D(\alpha^*)] - \frac{2L^2 \log(en)}{\lambda},$$

which yields

$$\mathbb{E}[D(\alpha^*) - D(\alpha^{(n)})] \leq \frac{2L^2 \log(en)}{\lambda n}.$$

7.4 Proof Of Theorem 11

The proof is identical to the proof of Theorem 2. We just need to notice that at the end of the first stage, we have $\mathbb{E}\epsilon_D^{(0)} \leq 2L^2 \log(en)/(\lambda n)$. It implies that $t_0 \leq \max(0, \lceil n \log(2\lambda n \cdot 2L^2 \log(en)/(\lambda n G)) \rceil)$.

7.5 Proof Of Proposition 15

Consider any feasible dual variable α and the corresponding $w = w(\alpha)$. Since

$$w = \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i x_i, \quad w^* = \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i^* x_i,$$

we have

$$\lambda(w - w^*)^\top w^* = \frac{1}{n} \sum_{i=1}^n (\alpha_i - \alpha_i^*) w^{*\top} x_i.$$

Therefore

$$\begin{aligned} & D(\alpha^*) - D(\alpha) \\ &= \frac{1}{n} \sum_{i=1}^n \left[\phi_i^*(-\alpha_i) - \phi_i^*(-\alpha_i^*) + (\alpha_i - \alpha_i^*) w^{*\top} x_i \right] + \frac{\lambda}{2} [w^\top w - w^{*\top} w^* - 2(w - w^*)^\top w^*] \\ &= \frac{1}{n} \sum_{i=1}^n \left[\phi_i^*(-\alpha_i) - \phi_i^*(-\alpha_i^*) + (\alpha_i - \alpha_i^*) w^{*\top} x_i \right] + \frac{\lambda}{2} (w - w^*)^\top (w - w^*). \end{aligned}$$

Since $w^{*\top} x_i \in \partial \phi_i^*(-\alpha_i^*)$, we have

$$\phi_i^*(-\alpha_i) - \phi_i^*(-\alpha_i^*) + (\alpha_i - \alpha_i^*) w^{*\top} x_i \geq \gamma_i (\alpha_i - \alpha_i^*)^2.$$

By combining the previous two displayed inequalities, we obtain the first desired bound.

Next, we let $u = w^{*\top} x_i$, $v = w^\top x_i$. Since $-a_i \in \partial \phi_i(v)$ and $-\alpha_i^* \in \partial \phi_i(u)$, it follows that $u \in \partial \phi_i^*(-\alpha_i^*)$ and $v \in \partial \phi_i^*(-a_i)$. Therefore

$$\begin{aligned} & |u - v| \cdot |\alpha_i^* - a_i| \\ &= \underbrace{[\phi_i^*(-a_i) - \phi_i^*(-\alpha_i^*) + u(a_i - \alpha_i^*)]}_{\geq 0} + \underbrace{[\phi_i^*(-\alpha_i^*) - \phi_i^*(-a_i) + v(\alpha_i^* - a_i)]}_{\geq 0} \\ &\geq \phi_i^*(-a_i) - \phi_i^*(-\alpha_i^*) + u(a_i - \alpha_i^*) \geq \gamma_i(u) |a_i - \alpha_i^*|^2. \end{aligned}$$

This implies the second bound.

7.6 Proof Of Theorem 16

The following lemma is very similar to Lemma 19 with nearly identical proof, but it focuses only on the convergence of dual objective function using (5).

Lemma 23 *Assume that (5) is valid. Then for any iteration t and any $s \in [0, 1]$ we have*

$$\mathbb{E}[D(\alpha^{(t)}) - D(\alpha^{(t-1)})] \geq \frac{s}{2n} \mathbb{E}[D(\alpha^*) - D(\alpha^{(t-1)})] + \frac{3s\lambda}{4n} \|w^* - w^{(t-1)}\|^2 - \left(\frac{s}{n}\right)^2 \frac{G_*^{(t)}(s)}{2\lambda},$$

where

$$G_*^{(t)}(s) = \frac{1}{n} \sum_{i=1}^n \left(\|x_i\|^2 - \frac{\gamma_i \lambda n}{s} \right) \mathbb{E}[(\alpha_i^* - \alpha_i^{(t-1)})^2].$$

Proof Since only the i 'th element of α is updated, the improvement in the dual objective can be written as

$$n[D(\alpha^{(t)}) - D(\alpha^{(t-1)})] = \underbrace{\left(-\phi^*(-\alpha_i^{(t)}) - \frac{\lambda n}{2} \|w^{(t)}\|^2 \right)}_{A_i} - \underbrace{\left(-\phi^*(-\alpha_i^{(t-1)}) - \frac{\lambda n}{2} \|w^{(t-1)}\|^2 \right)}_{B_i}.$$

By the definition of the update we have for all $s \in [0, 1]$ that

$$\begin{aligned} A_i &= \max_{\Delta \alpha_i} -\phi^*(-(\alpha_i^{(t-1)} + \Delta \alpha_i)) - \frac{\lambda n}{2} \|w^{(t-1)} + (\lambda n)^{-1} \Delta \alpha_i x_i\|^2 \\ &\geq -\phi^*(-(\alpha_i^{(t-1)} + s(\alpha_i^* - \alpha_i^{(t-1)}))) - \frac{\lambda n}{2} \|w^{(t-1)} + (\lambda n)^{-1} s(\alpha_i^* - \alpha_i^{(t-1)}) x_i\|^2. \end{aligned}$$

We can now apply the Jensen's inequality to obtain

$$\begin{aligned} A_i &\geq -s\phi_i^*(-\alpha_i^*) - (1-s)\phi_i^*(-\alpha_i^{(t-1)}) - \frac{\lambda n}{2} \|w^{(t-1)} + (\lambda n)^{-1} s(\alpha_i^* - \alpha_i^{(t-1)}) x_i\|^2 \\ &= -s[\phi_i^*(-\alpha_i^*) - \phi_i^*(-\alpha_i^{(t-1)})] - \underbrace{\phi_i^*(-\alpha_i^{(t-1)}) - \frac{\lambda n}{2} \|w^{(t-1)}\|^2}_{B_i} - s(\alpha_i^* - \alpha_i^{(t-1)}) x_i^\top w^{(t-1)} \\ &\quad - \frac{s^2(\alpha_i^* - \alpha_i^{(t-1)})^2}{2\lambda n} \|x_i\|^2. \end{aligned}$$

By summing over $i = 1, \dots, n$, we obtain

$$\begin{aligned} \sum_{i=1}^n [A_i - B_i] &\geq -s \sum_{i=1}^n [\phi_i^*(-\alpha_i^*) - \phi_i^*(-\alpha_i^{(t-1)})] - s \sum_{i=1}^n (\alpha_i^* - \alpha_i^{(t-1)}) x_i^\top w^{(t-1)} \\ &\quad - \frac{s^2}{2\lambda n} \sum_{i=1}^n (\alpha_i^* - \alpha_i^{(t-1)})^2 \|x_i\|^2 \\ &= -s \sum_{i=1}^n [\phi_i^*(-\alpha_i^*) - \phi_i^*(-\alpha_i^{(t-1)}) + \lambda(w^* - w^{(t-1)})^\top w^{(t-1)}] \\ &\quad - \frac{s^2}{2\lambda n} \sum_{i=1}^n (\alpha_i^* - \alpha_i^{(t-1)})^2 \|x_i\|^2, \end{aligned}$$

where the equality follows from $\sum_{i=1}^n (\alpha_i^* - \alpha_i^{(t-1)})x_i = \lambda n(w^* - w^{(t-1)})$. By rearranging the terms on the right hand side using $(w^* - w^{(t-1)})^\top w^{(t-1)} = \|w^*\|^2/2 - \|w^{(t-1)}\|^2/2 - \|w^* - w^{(t-1)}\|^2/2$, we obtain

$$\begin{aligned} & \sum_{i=1}^n [A_i - B_i] \\ & \geq -s \sum_{i=1}^n \left[\phi_i^*(-\alpha_i^*) - \phi_i^*(-\alpha_i^{(t-1)}) + \frac{\lambda}{2} \|w^*\|^2 - \frac{\lambda}{2} \|w^{(t-1)}\|^2 \right] + \frac{\lambda s n}{2} \|w^* - w^{(t-1)}\|^2 \\ & \quad - \frac{s^2}{2\lambda n} \sum_{i=1}^n (\alpha_i^* - \alpha_i^{(t-1)})^2 \|x_i\|^2 \\ & = s n [D(\alpha^*) - D(\alpha^{(t-1)})] + \frac{s \lambda n}{2} \|w^* - w^{(t-1)}\|^2 - \frac{s^2}{2\lambda n} \sum_{i=1}^n (\alpha_i^* - \alpha_i^{(t-1)})^2 \|x_i\|^2. \end{aligned}$$

We can now apply (5) to obtain

$$\begin{aligned} \sum_{i=1}^n [A_i - B_i] & \geq \frac{s n}{2} [D(\alpha^*) - D(\alpha^{(t-1)})] + \frac{3 s \lambda n}{4} \|w^* - w^{(t-1)}\|^2 \\ & \quad - \frac{s^2}{2\lambda n} \sum_{i=1}^n (\alpha_i^* - \alpha_i^{(t-1)})^2 (\|x_i\|^2 - \gamma_i \lambda n / s). \end{aligned}$$

This implies the desired result. \blacksquare

Lemma 24 Suppose that for all i , ϕ_i is L -Lipschitz. Let $G_*^{(t)}$ be as defined in Lemma 23. Then

$$G_*^{(t)}(s) \leq \frac{4L^2 N(s/(\lambda n))}{n}.$$

Proof Similarly to the proof of Lemma 22, we know that $(\alpha_i^* - \alpha_i^{(t-1)})^2 \leq 4L^2$. Moreover, $\|x_i\|^2 \leq 1$, and $\|x_i\|^2 - \frac{\gamma_i \lambda n}{s} \leq 0$ when $\gamma_i \geq s/(\lambda n)$. Therefore there are no more than $N(s/(\lambda n))$ data points i such that $\|x_i\|^2 - \frac{\gamma_i \lambda n}{s}$ is positive. The desired result follows from these facts. \blacksquare

Proof [Proof of Theorem 16] Let $\epsilon_D^{(t)} = \mathbb{E}[D(\alpha^*) - D(\alpha^{(t)})]$, and $G_*(s) = 4L^2 N(s/(\lambda n))/n$. We obtain from Lemma 23 and Lemma 24 that

$$\epsilon_D^{(t)} \leq (1 - s/(2n)) \epsilon_D^{(t-1)} + \left(\frac{s}{n}\right)^2 \frac{G_*(s)}{2\lambda}.$$

It follows that for all $t > 0$ we have

$$\begin{aligned} \epsilon_D^{(t)} & \leq (1 - s/(2n))^t \epsilon_D^{(0)} + \frac{1}{1 - (1 - s/(2n))} \left(\frac{s}{n}\right)^2 \frac{G_*(s)}{2\lambda} \\ & \leq e^{-st/2n} + \left(\frac{s}{n}\right) \frac{G_*(s)}{\lambda} \leq e^{-st/2n} + \epsilon_D/2. \end{aligned}$$

It follows that when

$$t \geq (2n/s) \log(2/\epsilon_D),$$

we have $\epsilon_D^{(t)} \leq \epsilon_D$. ■

7.7 Proof Of Theorem 18

Let $\epsilon_D^{(t)} = \mathbb{E}[D(\alpha^*) - D(\alpha^{(t)})]$. From Proposition 15, we know that for all $t \geq T_0$:

$$\begin{aligned} \epsilon_D^{(t)} &\geq \frac{1}{n} \sum_{i=1}^n \left[\gamma_i \mathbb{E} |\alpha_i^{(t)} - \alpha_i^*|^2 + \frac{\lambda}{2\rho} \mathbb{E} ((w^{(t)} - w^*)^\top x_i)^2 \right] \\ &\geq \frac{1}{n} \sum_{i=1}^n \left[\gamma_i \mathbb{E} |\alpha_i^{(t)} - \alpha_i^*|^2 + \frac{\lambda \gamma_i^2}{2\rho} \mathbb{E} (u_i^{(t-1)} - \alpha_i^*)^2 \right], \end{aligned}$$

where $-u_i^{(t-1)} \in \partial \phi_i(x_i^\top w^{(t)})$. It follows that given any $\gamma > 0$, we have

$$\begin{aligned} &\frac{1}{n} \sum_{i=1}^n \mathbb{E} |\alpha_i^{(t)} - u_i^{(t-1)}|^2 \\ &\leq \frac{N(\gamma)}{n} \sup_i \mathbb{E} |\alpha_i^{(t)} - u_i^{(t-1)}|^2 + \frac{2}{n} \sum_{i: \gamma_i \geq \gamma} \left[\mathbb{E} |\alpha_i^{(t)} - \alpha_i^*|^2 + \mathbb{E} (u_i^{(t-1)} - \alpha_i^*)^2 \right] \\ &\leq \frac{N(\gamma)}{n} \sup_i \mathbb{E} |\alpha_i^{(t)} - u_i^{(t-1)}|^2 + \frac{\frac{2}{n} \sum_{i=1}^n \left[\gamma_i \mathbb{E} |\alpha_i^{(t)} - \alpha_i^*|^2 + \frac{\lambda \gamma_i^2}{2\rho} \mathbb{E} (u_i^{(t-1)} - \alpha_i^*)^2 \right]}{\min(\gamma, \lambda \gamma^2 / (2\rho))} \\ &\leq \frac{N(\gamma)}{n} 4L^2 + \frac{2\epsilon_D^{(t)}}{\min(\gamma, \lambda \gamma^2 / (2\rho))}, \end{aligned}$$

where Lemma 22 is used for the last inequality. Since γ is arbitrary and $\epsilon_D^{(t)} \leq \epsilon_D$, it follows that

$$\frac{1}{n} \sum_{i=1}^n \mathbb{E} |\alpha_i^{(t)} - u_i^{(t-1)}|^2 \leq \tilde{\epsilon}_P.$$

Now plug into Lemma 19, we obtain for all $t \geq T_0 + 1$:

$$\begin{aligned} &\epsilon_D^{(t-1)} - \epsilon_D^{(t)} \\ &\geq \frac{s}{n} \mathbb{E}[P(w^{(t-1)}) - D(\alpha^{(t-1)})] - \left(\frac{s}{n}\right)^2 \frac{1}{2\lambda n} \sum_{i=1}^n \mathbb{E} [(u_i^{(t-1)} - \alpha_i^{(t-1)})^2] \\ &\geq \frac{s}{n} \mathbb{E}[P(w^{(t-1)}) - D(\alpha^{(t-1)})] - \left(\frac{s}{n}\right)^2 \frac{\tilde{\epsilon}_P}{2\lambda}. \end{aligned}$$

By taking $s = n/T_0$, and summing over $t = T_0 + 1, \dots, 2T_0 = T$, we obtain

$$\epsilon_D \geq \epsilon_D^{(T_0)} - \epsilon_D^{(T)} \geq \mathbb{E}[P(\bar{w}) - D(\bar{\alpha})] - \frac{\tilde{\epsilon}_P}{2\lambda T_0}.$$

This proves the desired bound.

8. Experimental Results

In this section we demonstrate the tightness of our theory. All our experiments are performed with the smooth variant of the hinge-loss defined in (6), where the value of γ is taken from the set $\{0, 0.01, 0.1, 1\}$. Note that for $\gamma = 0$ we obtain the vanilla non-smooth hinge-loss.

In the experiments, we use ϵ_D to denote the dual sub-optimality, and ϵ_P to denote the primal sub-optimality (note that this is different than the notation in our analysis which uses ϵ_P to denote the duality gap). It follows that $\epsilon_D + \epsilon_P$ is the duality gap.

8.1 Data

The experiments were performed on three large data sets with very different feature counts and sparsity, which were kindly provided by Thorsten Joachims. The astro-ph data set classifies abstracts of papers from the physics ArXiv according to whether they belong in the astro-physics section; CCAT is a classification task taken from the Reuters RCV1 collection; and cov1 is class 1 of the covtype data set of Blackard, Jock & Dean. The following table provides details of the data set characteristics.

Data Set	Training Size	Testing Size	Features	Sparsity
astro-ph	29882	32487	99757	0.08%
CCAT	781265	23149	47236	0.16%
cov1	522911	58101	54	22.22%

8.2 Linear Convergence For Smooth Hinge-loss

Our first experiments are with ϕ_γ where we set $\gamma = 1$. The goal of the experiment is to show that the convergence is indeed linear. We ran the SDCA algorithm for solving the regularized loss minimization problem with different values of regularization parameter λ . Figure 1 shows the results. Note that a logarithmic scale is used for the vertical axis. Therefore, a straight line corresponds to linear convergence. We indeed observe linear convergence for the duality gap.

8.3 Convergence For Non-smooth Hinge-loss

Next we experiment with the original hinge loss, which is 1-Lipschitz but is not smooth. We again ran the SDCA algorithm for solving the regularized loss minimization problem with different values of regularization parameter λ . Figure 2 shows the results. As expected, the overall convergence rate is slower than the case of a smoothed hinge-loss. However, it is also apparent that for large values of λ a linear convergence is still exhibited, as expected according to our refined analysis. The bounds plotted are based on Theorem 2, which are slower than what we observe, as expected from the refined analysis in Section 5.

8.4 Effect Of Smoothness Parameter

We next show the effect of the smoothness parameter. Figure 3 shows the effect of the smoothness parameter on the rate of convergence. As can be seen, the convergence becomes faster as the loss function becomes smoother. However, the difference is more dominant when λ decreases.

Figure 4 shows the effect of the smoothness parameter on the zero-one test error. It is noticeable that even though the non-smooth hinge-loss is considered a tighter approximation of the zero-one error, in most cases, the smoothed hinge-loss actually provides a lower test error than the non-smooth hinge-loss. In any case, it is apparent that the smooth hinge-loss decreases the zero-one test error faster than the non-smooth hinge-loss.

8.5 Cyclic vs. Stochastic vs. Random Permutation

In Figure 5 we compare choosing dual variables at random with repetitions (as done in SDCA) vs. choosing dual variables using a random permutation at each epoch (as done in SDCA-Perm) vs. choosing dual variables in a fixed cyclic order (that was chosen once at random). As can be seen, a cyclic order does not lead to linear convergence and yields actual convergence rate much slower than the other methods and even worse than our bound. As mentioned before, some of the earlier analyses such as Luo and Tseng (1992) can be applied both to stochastic and to cyclic dual coordinate ascent methods with similar results. This means that their analysis, which can be no better than the behavior of cyclic dual coordinate ascent, is inferior to our analysis. Finally, we also observe that SDCA-Perm is sometimes faster than SDCA.

8.6 Comparison To SGD

We next compare SDCA to Stochastic Gradient Descent (SGD). In particular, we implemented SGD with the update rule $w^{(t+1)} = (1 - 1/t)w^{(t)} - \frac{1}{\lambda t} \phi'_i(w^{(t)\top} x_i) x_i$, where i is chosen uniformly at random and ϕ'_i denotes a sub-gradient of ϕ_i . One clear advantage of SDCA is the availability of a clear stopping condition (by calculating the duality gap). In Figure 6 and Figure 7 we present the primal sub-optimality of SDCA, SDCA-Perm, and SGD. As can be seen, SDCA converges faster than SGD in most regimes. SGD can be better if both λ is high and one performs a very small number of epochs. This is in line with our theory of Section 4. However, SDCA quickly catches up.

In Figure 8 we compare the zero-one test error of SDCA, when working with the smooth hinge-loss ($\gamma = 1$) to the zero-one test error of SGD, when working with the non-smooth hinge-loss. As can be seen, SDCA with the smooth hinge-loss achieves the smallest zero-one test error faster than SGD.

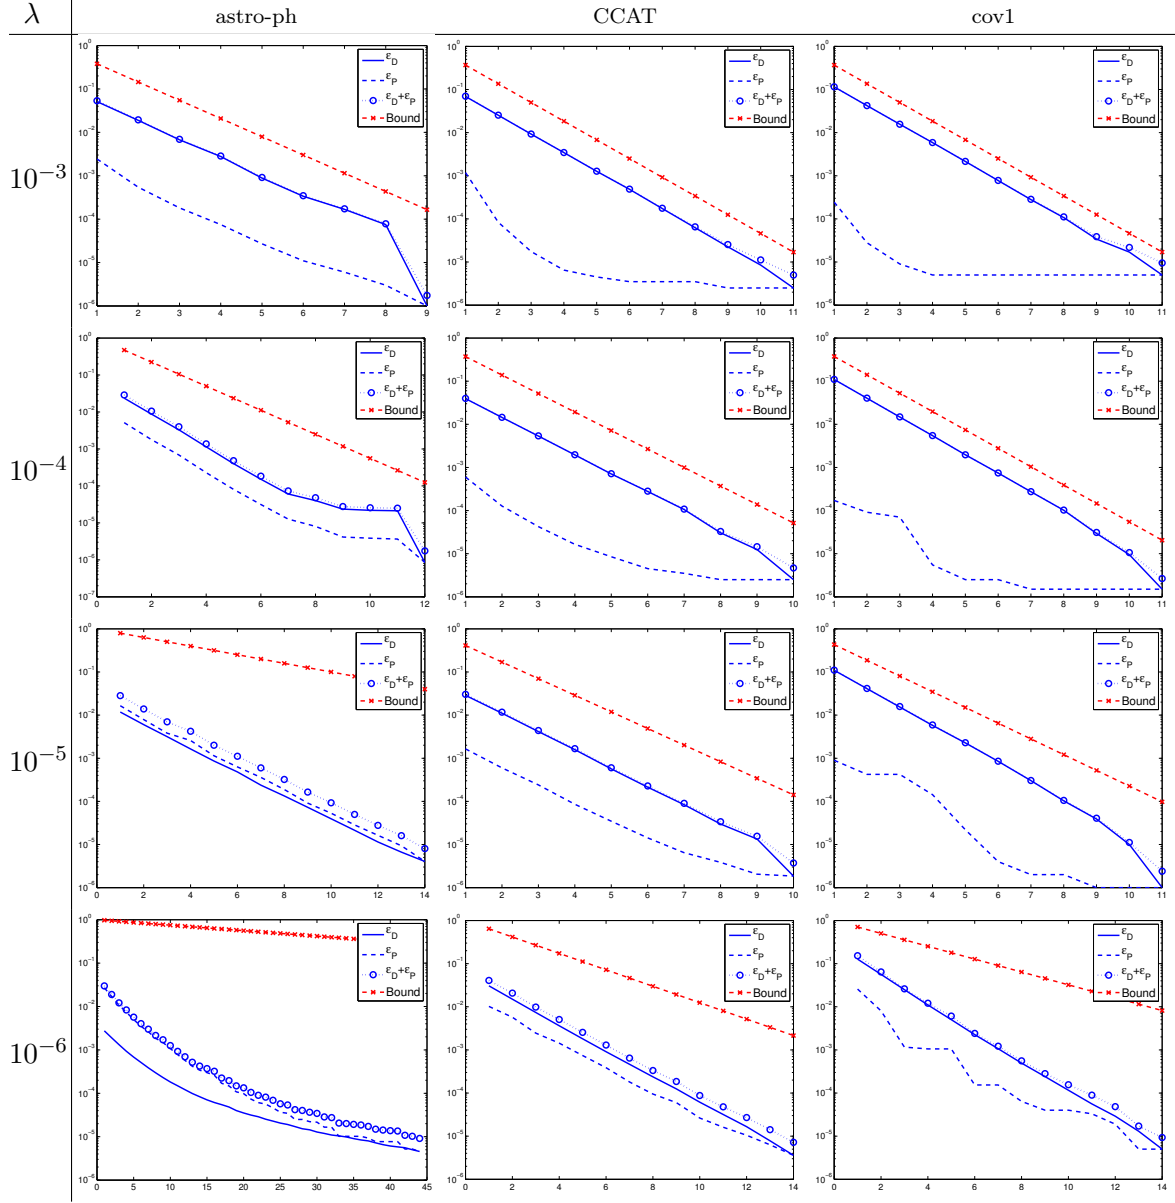


Figure 1: Experiments with the smoothed hinge-loss ($\gamma = 1$). The primal and dual sub-optimality, the duality gap, and our bound are depicted as a function of the number of epochs, on the astro-ph (left), CCAT (center) and cov1 (right) data sets. In all plots the horizontal axis is the number of iterations divided by training set size (corresponding to the number of epochs through the data).

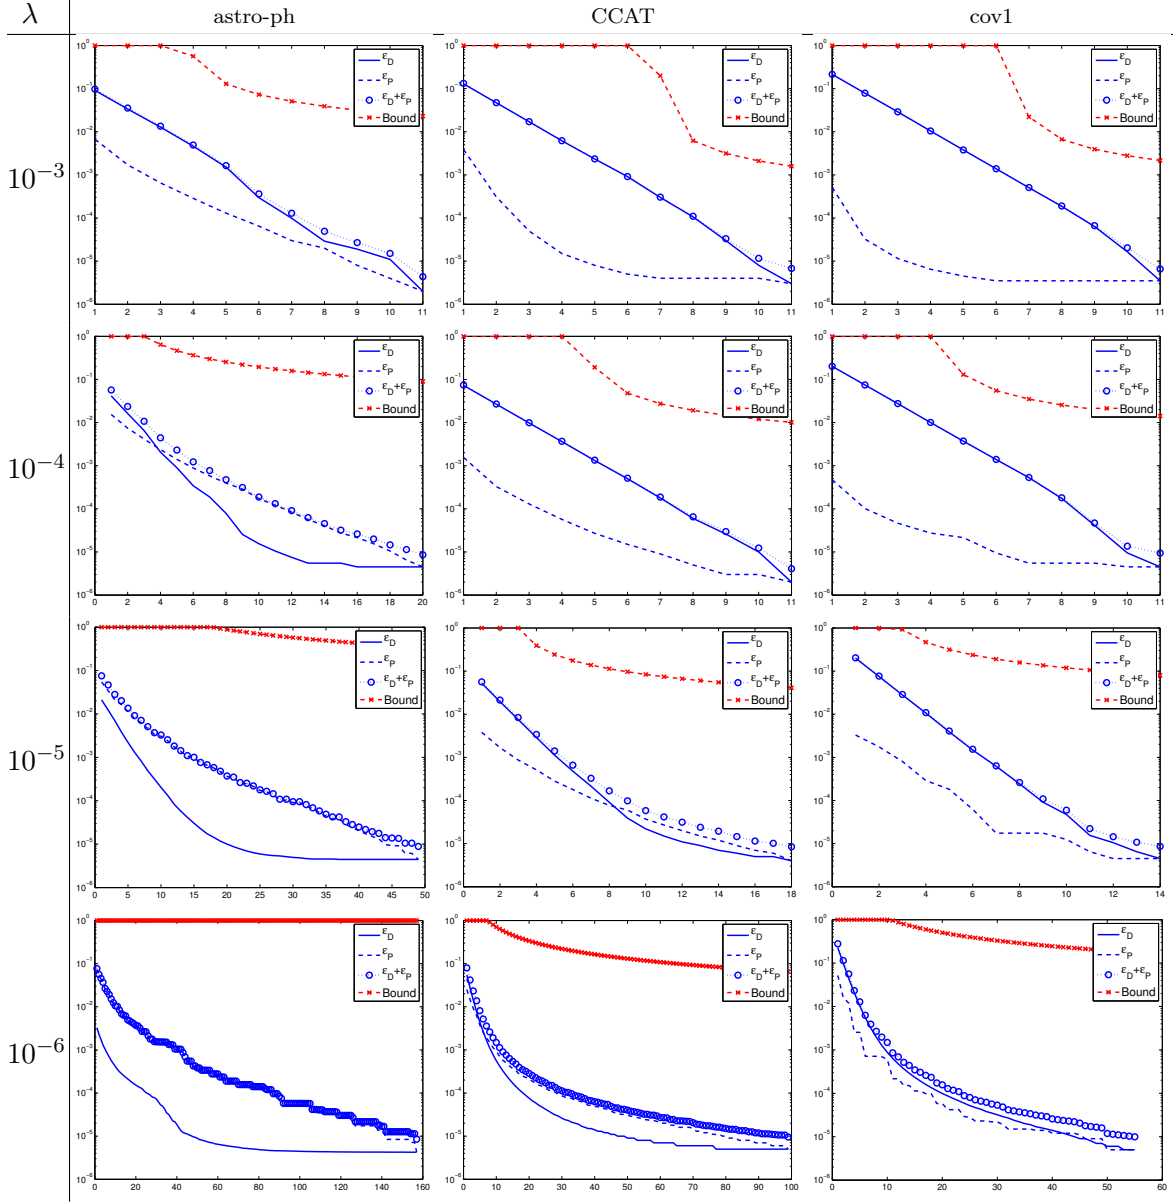
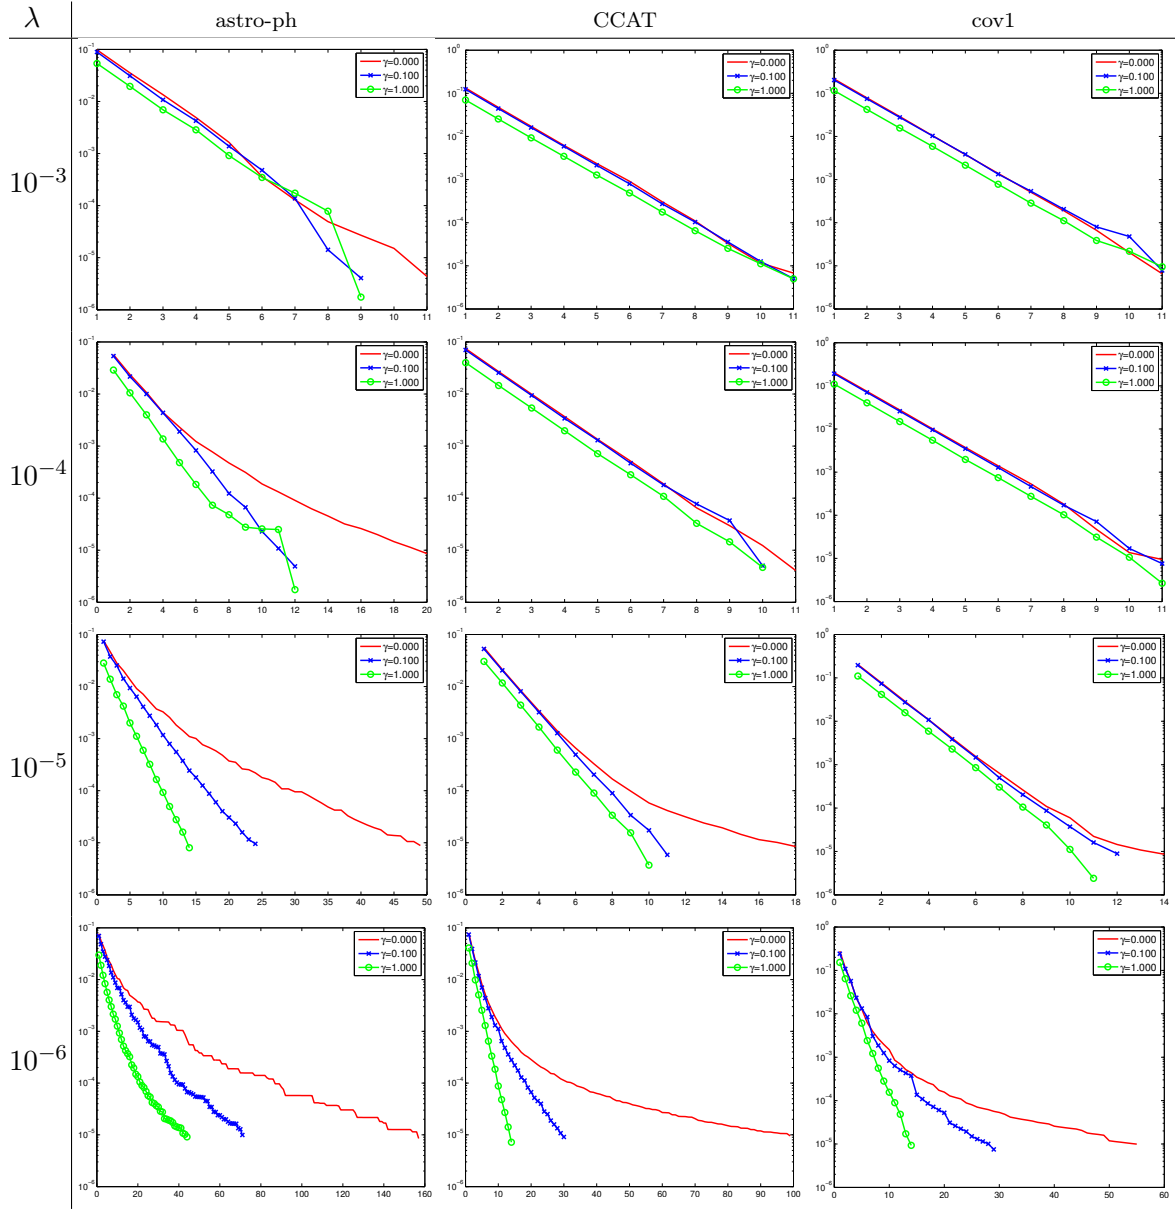


Figure 2: Experiments with the hinge-loss (non-smooth). The primal and dual suboptimality, the duality gap, and our bound are depicted as a function of the number of epochs, on the astro-ph (left), CCAT (center) and cov1 (right) data sets. In all plots the horizontal axis is the number of iterations divided by training set size (corresponding to the number of epochs through the data).


 Figure 3: Duality gap as a function of the number of rounds for different values of γ .

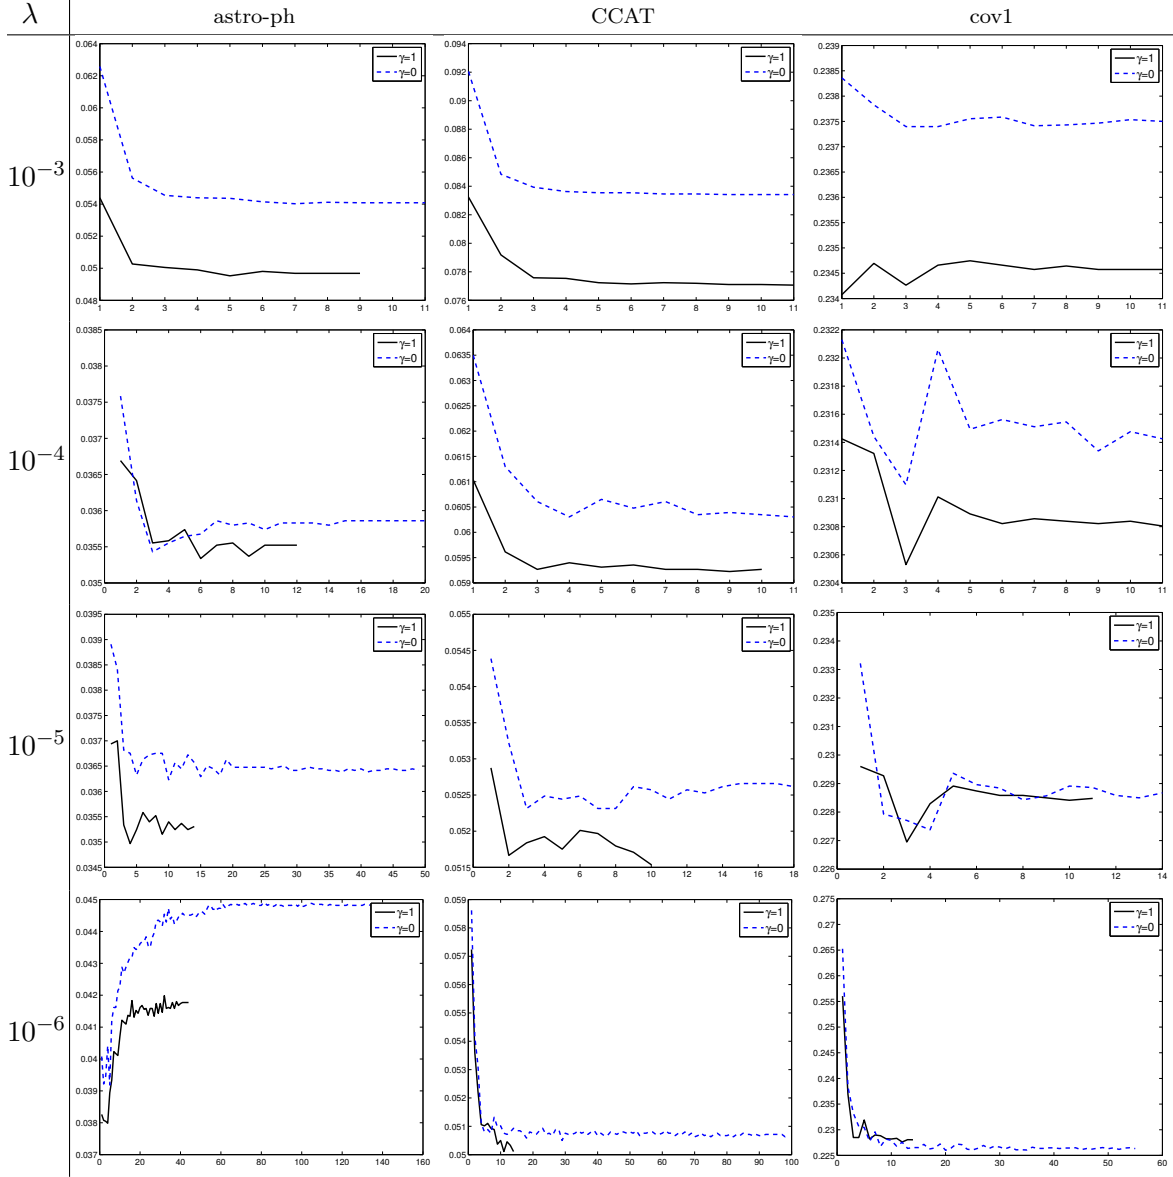


Figure 4: Comparing the test zero-one error of SDCA for smoothed hinge-loss ($\gamma = 1$) and non-smooth hinge-loss ($\gamma = 0$). In all plots the vertical axis is the zero-one error on the test set and the horizontal axis is the number of iterations divided by training set size (corresponding to the number of epochs through the data). We terminated each method when the duality gap was smaller than 10^{-5} .

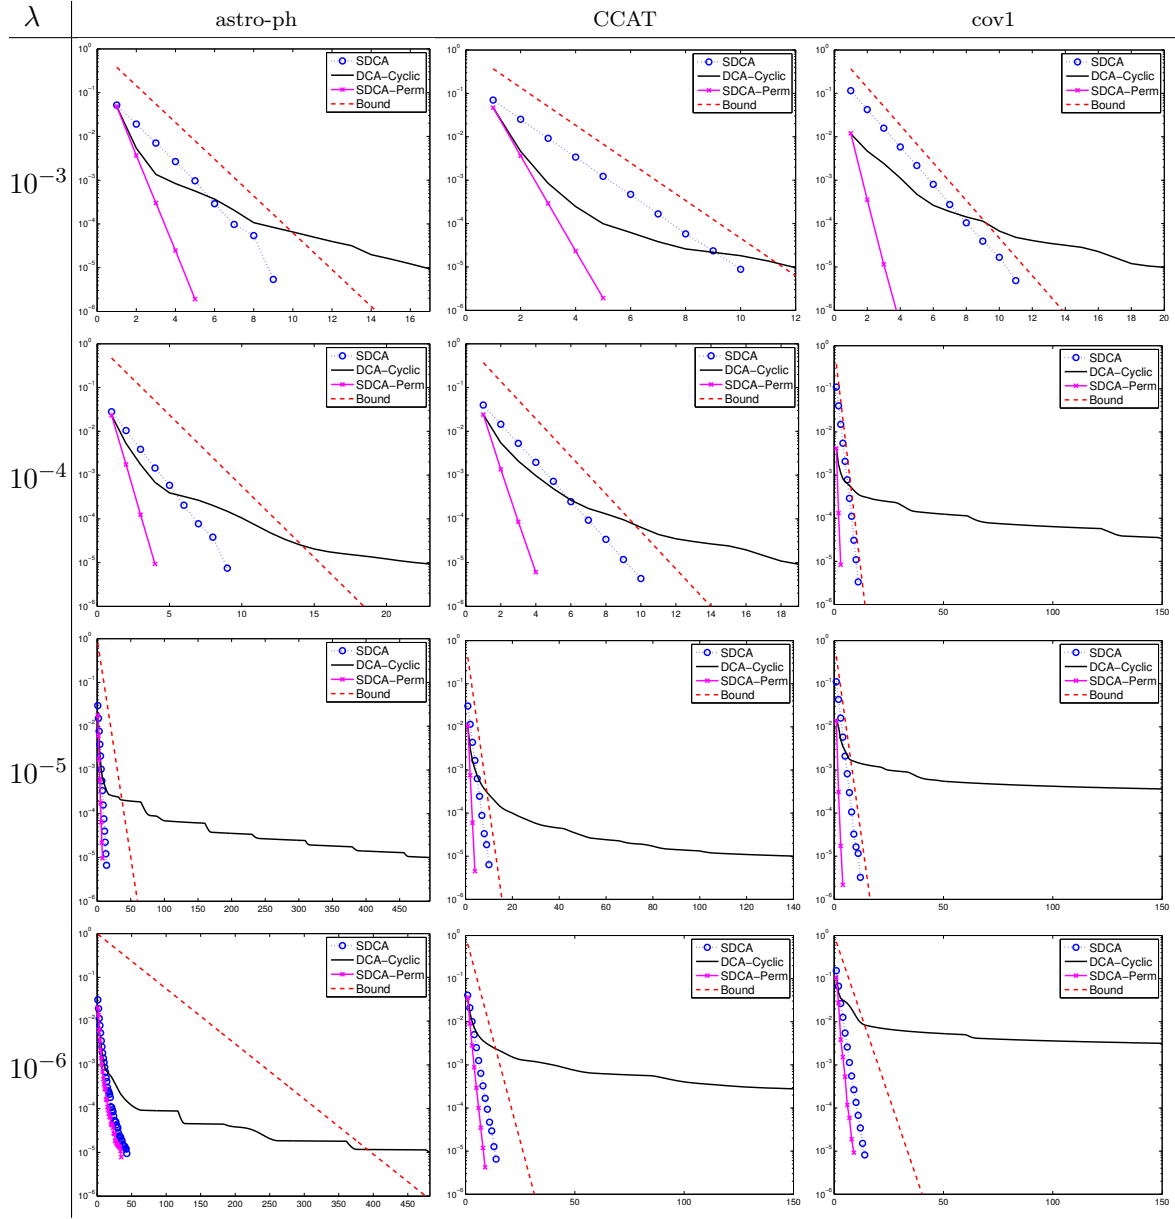


Figure 5: Comparing the duality gap achieved by choosing dual variables at random with repetitions (SDCA), choosing dual variables at random without repetitions (SDCA-Perm), or using a fixed cyclic order. In all cases, the duality gap is depicted as a function of the number of epochs for different values of λ . The loss function is the smooth hinge loss with $\gamma = 1$.

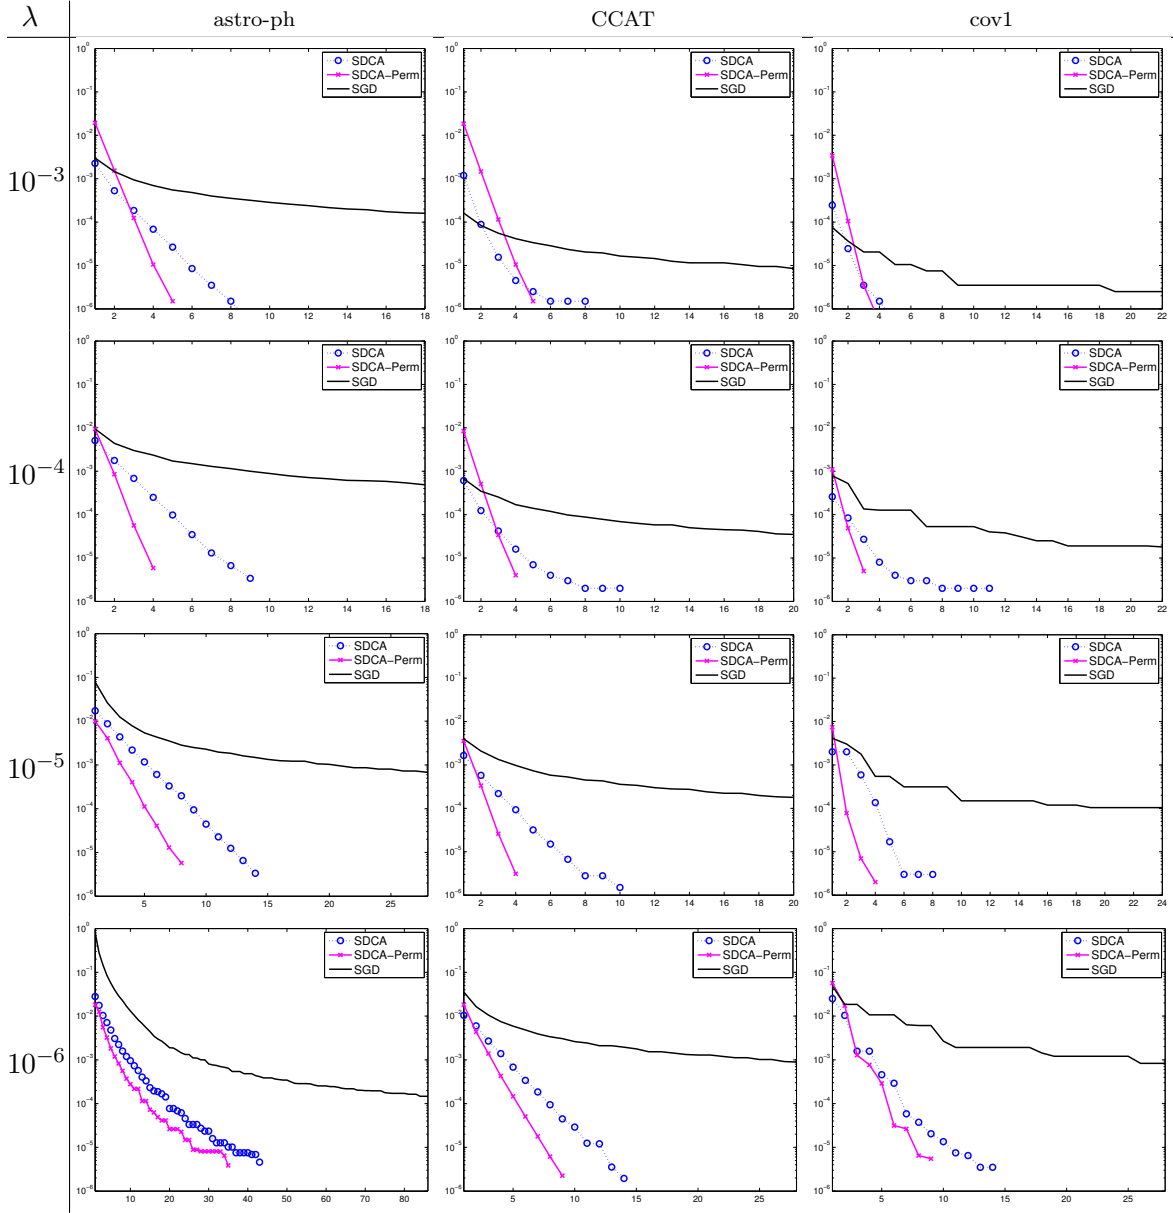


Figure 6: Comparing the primal sub-optimality of SDCA and SGD for the smoothed hinge-loss ($\gamma = 1$). In all plots the horizontal axis is the number of iterations divided by training set size (corresponding to the number of epochs through the data).

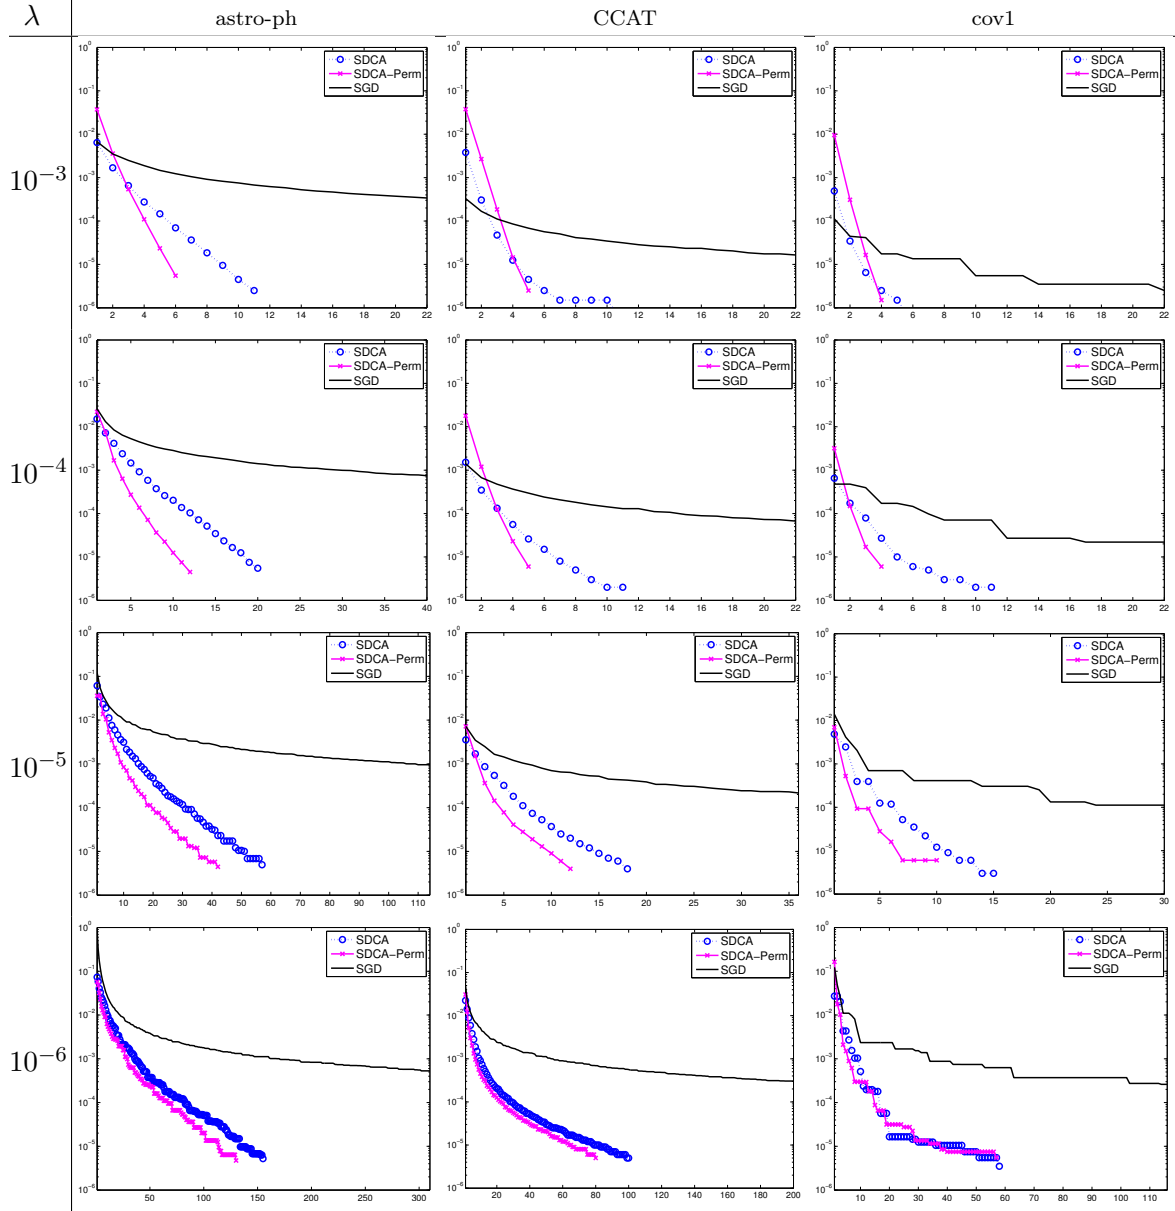


Figure 7: Comparing the primal sub-optimality of SDCA and SGD for the non-smooth hinge-loss ($\gamma = 0$). In all plots the horizontal axis is the number of iterations divided by training set size (corresponding to the number of epochs through the data).

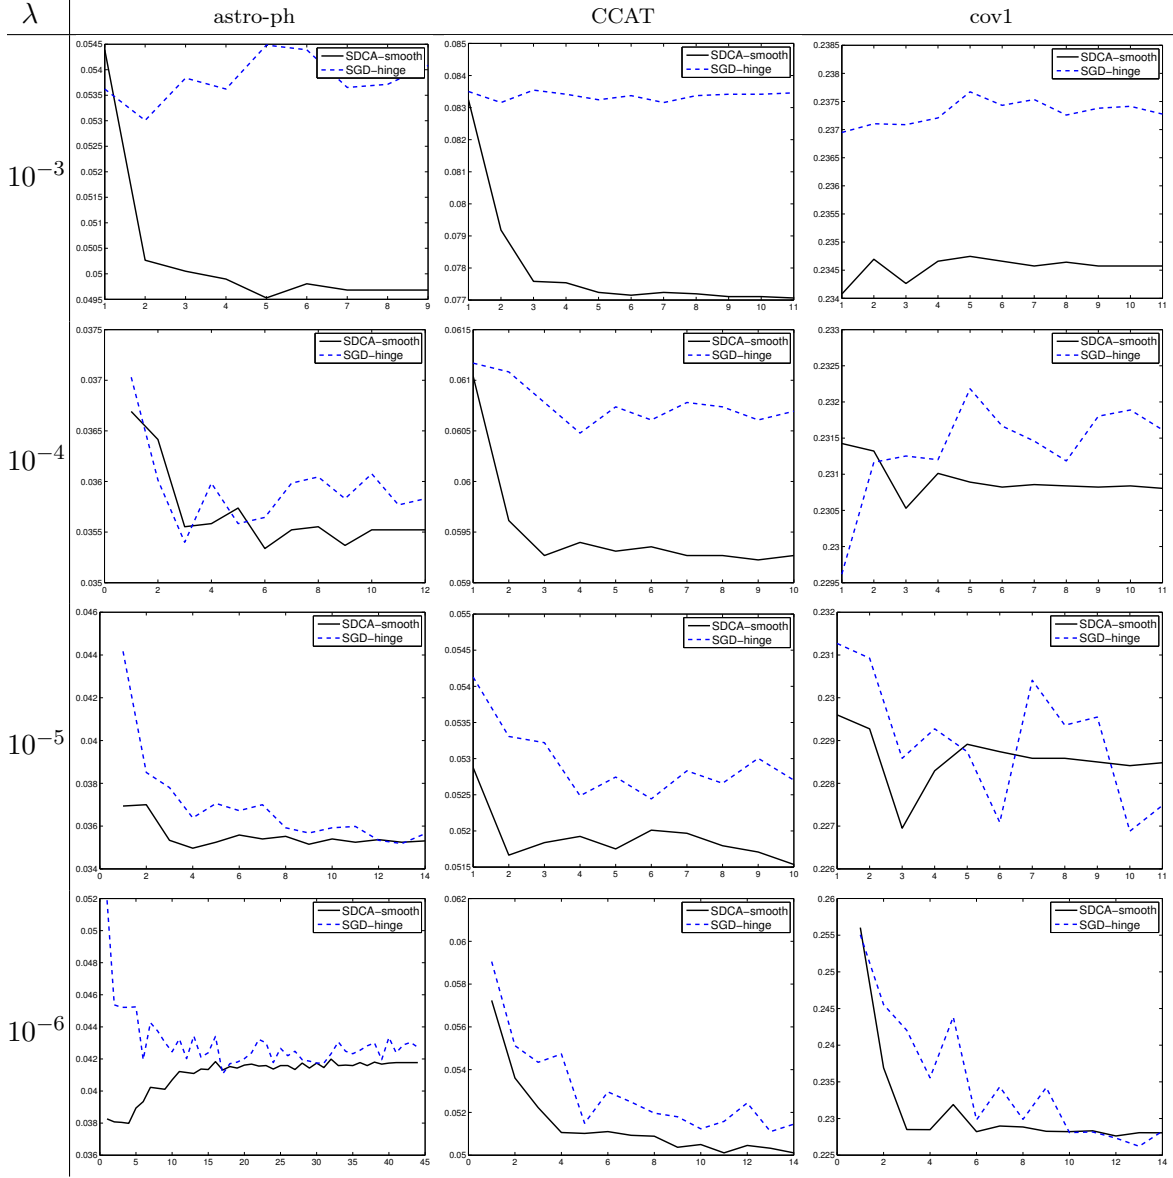


Figure 8: Comparing the test error of SDCA with the smoothed hinge-loss ($\gamma = 1$) to the test error of SGD with the non-smoothed hinge-loss. In all plots the vertical axis is the zero-one error on the test set and the horizontal axis is the number of iterations divided by training set size (corresponding to the number of epochs through the data). We terminated SDCA when the duality gap was smaller than 10^{-5} .

Acknowledgments

Shai Shalev-Shwartz acknowledges the support of the Israeli Science Foundation grant number 598-10. Tong Zhang acknowledges the support of NSF grant DMS-1007527 and NSF grant IIS-1016061.

References

- L. Bottou and O. Bousquet. The tradeoffs of large scale learning. In *Advances in Neural Information Processing Systems (NIPS)*, pages 161–168, 2008.
- M. Collins, A. Globerson, T. Koo, X. Carreras, and P. Bartlett. Exponentiated gradient algorithms for conditional random fields and max-margin markov networks. *Journal of Machine Learning Research*, 9:1775–1822, 2008.
- Y. Le Cun and L. Bottou. Large scale online learning. In *Advances in Neural Information Processing Systems 16: Proceedings of the 2003 Conference*, volume 16, page 217. MIT Press, 2004.
- C.J. Hsieh, K.W. Chang, C.J. Lin, S.S. Keerthi, and S. Sundararajan. A dual coordinate descent method for large-scale linear SVM. In *Proceedings of the International Conference on Machine Learning (ICML)*, pages 408–415, 2008.
- D. Hush, P. Kelly, C. Scovel, and I. Steinwart. QP algorithms with guaranteed accuracy and run time for support vector machines. *Journal of Machine Learning Research*, 7:733–769, 2006.
- T. Joachims. Making large-scale support vector machine learning practical. In B. Schölkopf, C. Burges, and A. Smola, editors, *Advances in Kernel Methods - Support Vector Learning*. MIT Press, 1998.
- S. Lacoste-Julien, M. Jaggi, M. Schmidt, and P. Pletscher. Stochastic block-coordinate frank-wolfe optimization for structural svms. *arXiv preprint:1207.4747*, 2012.
- N. Le Roux, M. Schmidt, and F. Bach. A Stochastic Gradient Method with an Exponential Convergence Rate for Strongly-Convex Optimization with Finite Training Sets. In *Advances in Neural Information Processing Systems (NIPS): Proceedings of the 2012 Conference*, 2012. arXiv preprint:1202.6258.
- Z.Q. Luo and P. Tseng. On the convergence of coordinate descent method for convex differentiable minimization. *J. Optim. Theory Appl.*, 72:7–35, 1992.
- O. Mangasarian and D. Musicant. Successive overrelaxation for support vector machines. *IEEE Transactions on Neural Networks*, 10, 1999.
- N. Murata. A statistical study of on-line learning. *Online Learning and Neural Networks*. Cambridge University Press, Cambridge, UK, 1998.
- Y. Nesterov. Efficiency of coordinate descent methods on huge-scale optimization problems. *SIAM Journal on Optimization*, 22(2):341–362, 2012.

- J. C. Platt. Fast training of Support Vector Machines using sequential minimal optimization. In B. Schölkopf, C. Burges, and A. Smola, editors, *Advances in Kernel Methods - Support Vector Learning*. MIT Press, 1998.
- H. Robbins and S. Monro. A stochastic approximation method. *The Annals of Mathematical Statistics*, pages 400–407, 1951.
- S. Shalev-Shwartz and N. Srebro. SVM optimization: Inverse dependence on training set size. In *International Conference on Machine Learning (ICML)*, pages 928–935, 2008.
- S. Shalev-Shwartz and A. Tewari. Stochastic methods for l_1 regularized loss minimization. In *Proceedings of the International Conference on Machine Learning (ICML)*, page 117, 2009.
- S. Shalev-Shwartz, Y. Singer, and N. Srebro. Pegasos: Primal Estimated sub-GrAdient SOLver for SVM. In *Proceedings of the International Conference on Machine Learning (ICML)*, pages 807–814, 2007.
- K. Sridharan, N. Srebro, and S. Shalev-Shwartz. Fast rates for regularized objectives. In *Advances in Neural Information Processing Systems (NIPS)*, 2009.
- T. Zhang. Solving large scale linear prediction problems using stochastic gradient descent algorithms. In *Proceedings of the Twenty-First International Conference on Machine Learning (ICML)*, 2004.